

ขั้นตอนวิธีกราฟ Graph Algorithms

การอบรมคอมพิวเตอร์โอลิมปิก สอวน. ค่ายที่ 2 ปีการศึกษา 2563



อ.ดร.สุภาปนา บุญชู

สาขาวิชาวิทยาการคอมพิวเตอร์ คณะวิทยาศาสตร์และเทคโนโลยี

มหาวิทยาลัยธรรมศาสตร์

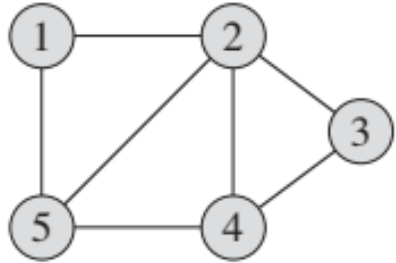
Graph Fundamentals

Representing a Graph

- We can choose between two standard ways to represent a graph $G = (V, E)$ as a collection of **adjacency lists** or as an **adjacency matrix**.
- Because the adjacency-list representation provides a compact way to represent **sparse** graphs—those for which $|E|$ is much less than $|V|^2$ —it is usually the method of choice.
- We may prefer an adjacency-matrix representation, however, when the graph is **dense**— $|E|$ is close to $|V|^2$ —or when we need to be able to tell quickly if there is an edge connecting two given vertices.

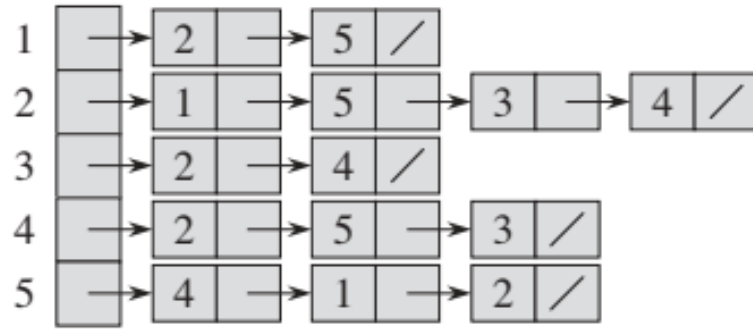
Representing a Graph

Graph



(a)

Represented using Adjacency Lists.



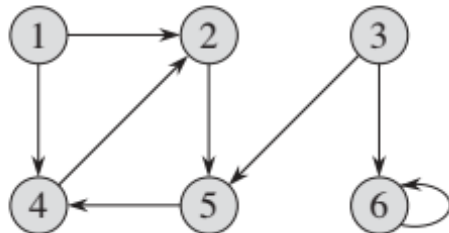
(b)

Represented using Adjacency Matrix.

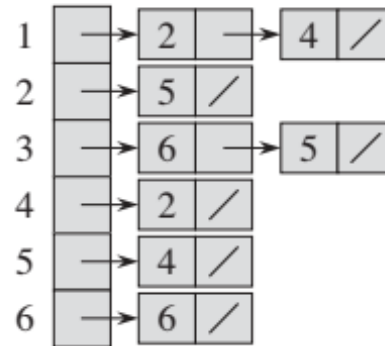
	1	2	3	4	5
1	0	1	0	0	1
2	1	0	1	1	1
3	0	1	0	1	0
4	0	1	1	0	1
5	1	1	0	1	0

(c)

Undirected graph



(a)



(b)

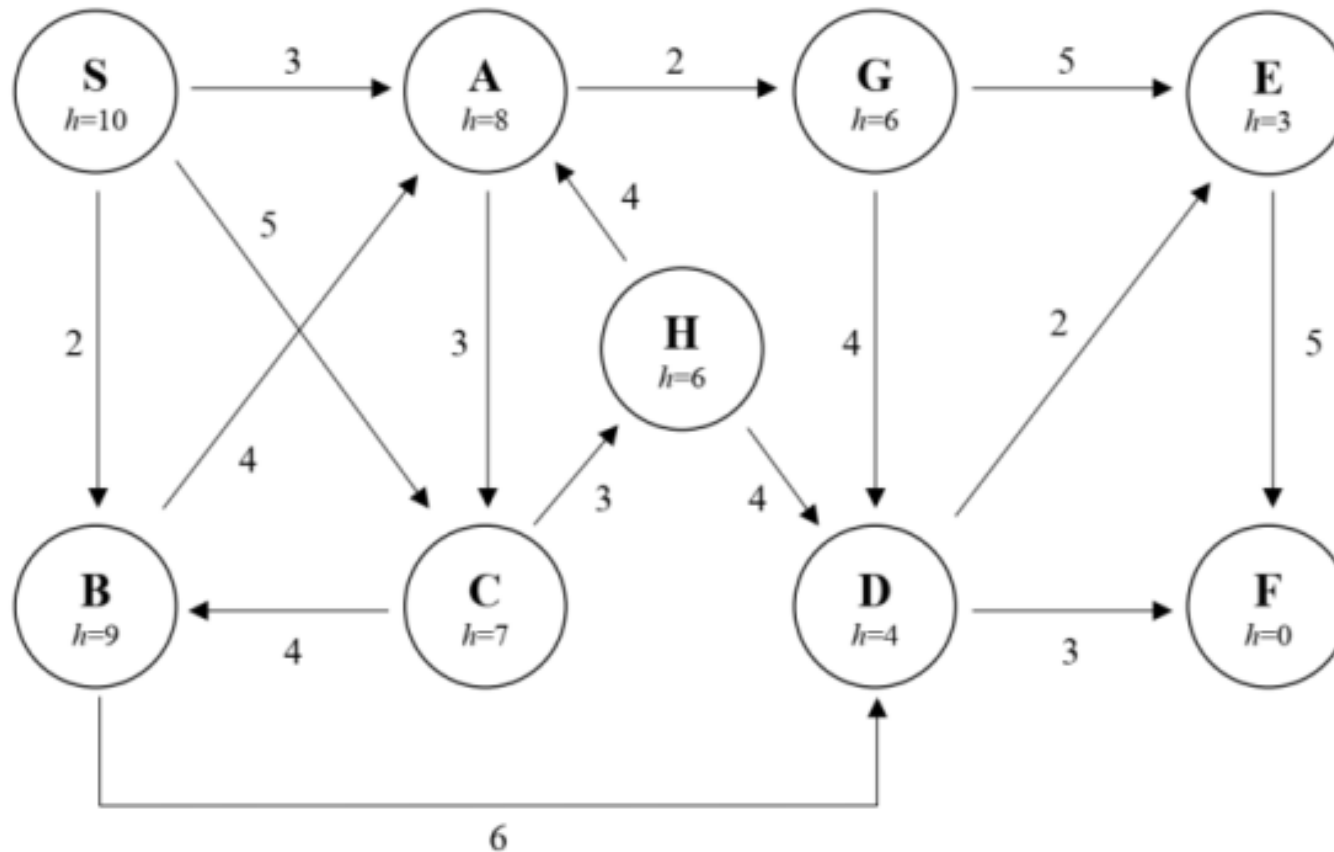
	1	2	3	4	5	6
1	0	1	0	1	0	0
2	0	0	0	0	1	0
3	0	0	0	0	1	1
4	0	1	0	0	0	0
5	0	0	0	1	0	0
6	0	0	0	0	0	1

(c)

Directed graph

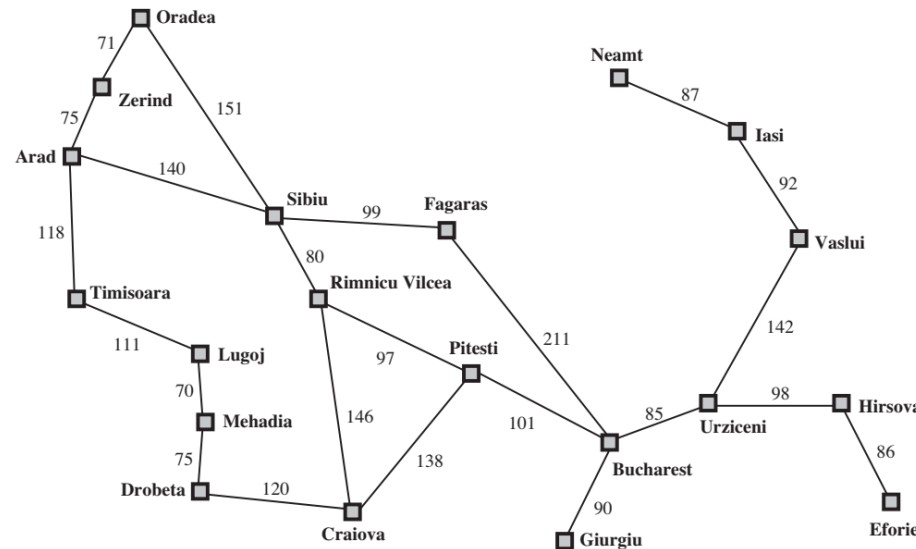
Exercise 1

- ให้นักเรียนเขียน โปรแกรมเพื่อ แทนค่า กราฟต่อไปนี้ด้วยวิธี Adjacency Matrix



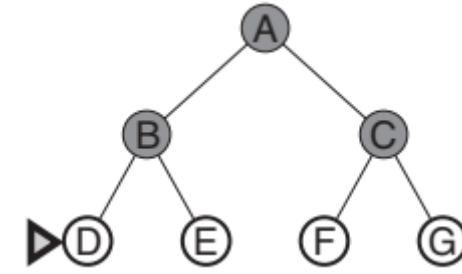
Problem of Graph Searching

- The problem is to find the solution for the **goal** on a given graph.
- A **solution** is an action sequence, so search algorithms work by considering various possible action sequences.
- **Solution quality** is measured by the path cost function, and an **optimal solution** has the lowest path cost among all solutions.



Uninformed Search Strategies (Blind Search)

Breadth-First Search (BFS)

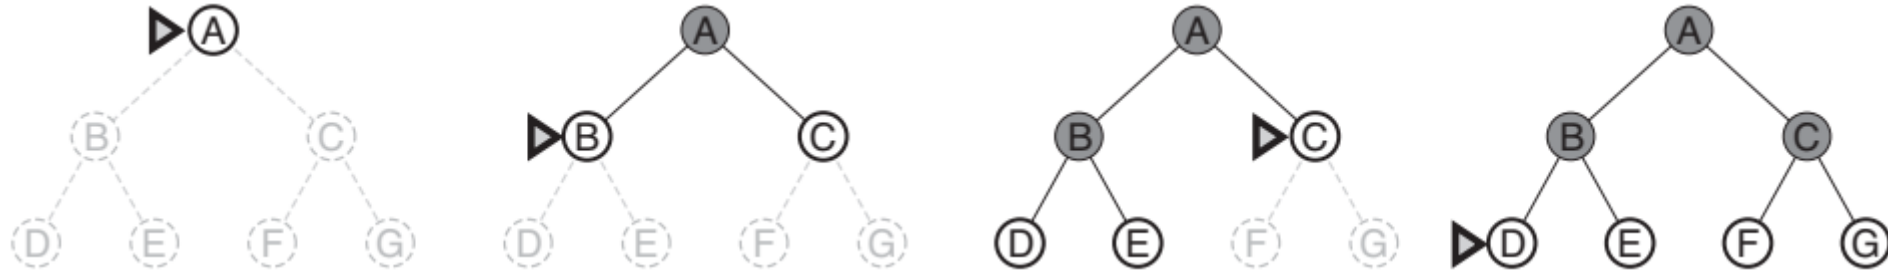


- **Root** is expanded first.
- Then, all the *successors* of the root node are expanded next, then *their successors*, and so on.
- In general, all the nodes are expanded at a **given depth** in the search tree **before** any nodes at the next level are expanded.
- BFS is achieved very simply by using a **FIFO queue** for the *frontier*.
 - Thus, new nodes (which are always deeper than their parents) go to the back of the queue, and **old nodes**, which are shallower than the new nodes, **get expanded first**.
- The algorithm, following the general template for graph search, discards any new path to a state already in the frontier or explored set; it is easy to see that any such path must be at least as deep as the one already found.
- There is one **slight tweak** on the general graph-search algorithm, which is that **the goal test is applied to each node when it is generated** rather than when it is selected for expansion.
 - Thus, BFS always has the shallowest path to every node on the frontier.

Breadth-First Search (BFS)

- Four criteria:
 - BFS is *complete*.
 - As soon as a **goal** node is generated, we know it is the **shallowest goal node** because all shallower nodes must have been generated already and failed the goal test.
 - Now, the **shallowest goal node** is not *necessarily* the *optimal* one; technically, breadth-first search is *optimal* if the path cost is a **nondecreasing function of the depth** of the node.
 - Complexity, both space and time complexity are $O(b^d)$.
- The memory requirements are a bigger problem for breadth-first search than is the execution time.
 - One might wait 13 days for the solution to an important problem with search depth 12, but no personal computer has the petabyte of memory it would take.
- In general, **exponential-complexity** search problems cannot be solved by uninformed methods for any but the smallest instances.

BFS



Depth	Nodes	Time	Memory
2	110	.11 milliseconds	107 kilobytes
4	11,110	11 milliseconds	10.6 megabytes
6	10^6	1.1 seconds	1 gigabyte
8	10^8	2 minutes	103 gigabytes
10	10^{10}	3 hours	10 terabytes
12	10^{12}	13 days	1 petabyte
14	10^{14}	3.5 years	99 petabytes
16	10^{16}	350 years	10 exabytes

Time and memory requirements for breadth-first search. The numbers shown assume branching factor $b = 10$; 1 million nodes/second; 1000 bytes/node.

Breadth-First Search (BFS)

```
function BREADTH-FIRST-SEARCH(problem) returns a solution, or failure  
  node  $\leftarrow$  a node with STATE = problem.INITIAL-STATE, PATH-COST = 0  
  if problem.GOAL-TEST(node.STATE) then return SOLUTION(node)  
  frontier  $\leftarrow$  a FIFO queue with node as the only element  
  explored  $\leftarrow$  an empty set  
  loop do  
    if EMPTY?(frontier) then return failure  
    node  $\leftarrow$  POP(frontier) /* chooses the shallowest node in frontier */  
    add node.STATE to explored  
    for each action in problem.ACTIONS(node.STATE) do  
      child  $\leftarrow$  CHILD-NODE(problem, node, action)  
      if child.STATE is not in explored or frontier then  
        if problem.GOAL-TEST(child.STATE) then return SOLUTION(child)  
        frontier  $\leftarrow$  INSERT(child, frontier)
```

Uniform-cost search

- When **all step costs are equal**, breadth-first search is *optimal* because it always expands the *shallowest unexpanded* node.
- By a simple extension, we can find an algorithm that is optimal with any **step-cost** function.
 - Instead of expanding the shallowest node, uniform-cost search expands the node n with the lowest path cost $g(n)$. This is done by storing the frontier as a priority queue ordered by g .
 - Two differences from BFS:
 - The first is that the goal test is applied to a node when it is selected for expansion rather than when it is first generated.
 - The reason is that the first goal node that is generated may be on a suboptimal path.
 - The second difference is that a test is added in case a better path is found to a node currently on the frontier.

Uniform-cost search

function UNIFORM-COST-SEARCH(*problem*) **returns** a solution, or failure

node \leftarrow a node with STATE = *problem*.INITIAL-STATE, PATH-COST = 0

frontier \leftarrow a priority queue ordered by PATH-COST, with *node* as the only element

explored \leftarrow an empty set

loop do

if EMPTY?(*frontier*) **then return** failure

node \leftarrow POP(*frontier*) /* chooses the lowest-cost node in *frontier* */

if *problem*.GOAL-TEST(*node*.STATE) **then return** SOLUTION(*node*)

 add *node*.STATE to *explored*

for each *action* **in** *problem*.ACTIONS(*node*.STATE) **do**

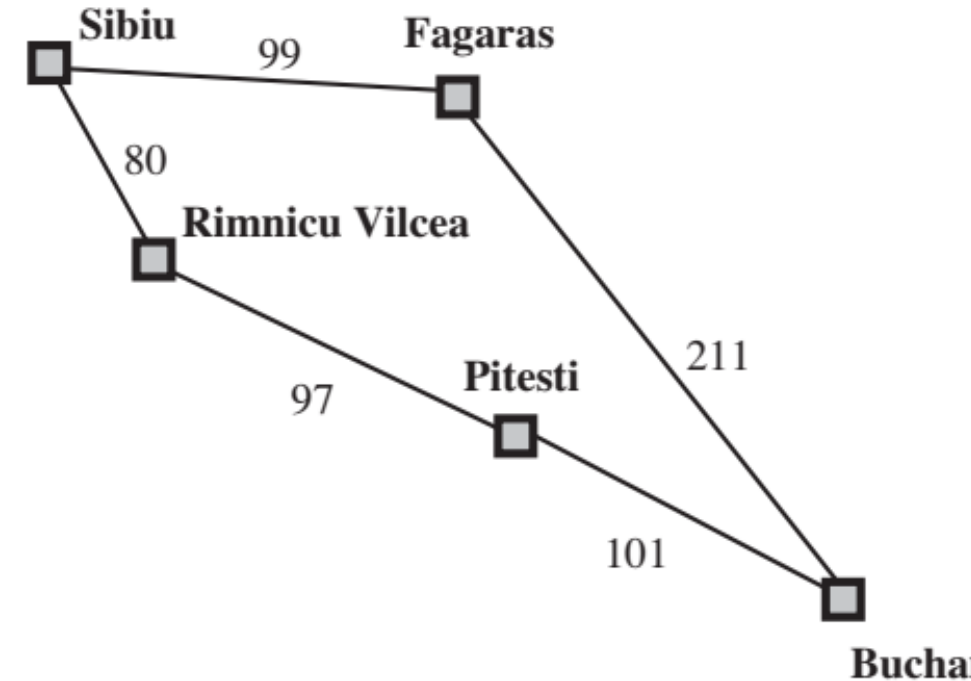
child \leftarrow CHILD-NODE(*problem*, *node*, *action*)

if *child*.STATE is not in *explored* or *frontier* **then**

frontier \leftarrow INSERT(*child*, *frontier*)

else if *child*.STATE is in *frontier* with higher PATH-COST **then**

 replace that *frontier* node with *child*



Uniform-cost search on a graph. The algorithm is identical to the general graph search algorithm, except for the use of a priority queue and the addition of an extra check in case a shorter path to a frontier state is discovered. The data structure for frontier needs to support efficient membership testing, so it should combine the capabilities of a priority queue and a hash table.

Depth-First Search (DFS)

- Depth-first search always expands the **deepest node** in the current frontier of the search tree.
 - The search proceeds immediately to the deepest level of the search tree, where **the nodes have no successors**.
 - Whereas breadth-first-search uses a **FIFO** queue, depth-first search uses a **LIFO** queue.
 - A LIFO queue means that the **most recently generated node** is chosen for expansion.
- It is common to implement depth-first search with a **recursive** function that calls itself on each of its children in turn.
- The graph-search version, which avoids repeated states and redundant paths, is **complete** in finite state spaces because it will eventually expand every node.
- The tree-search version, on the other hand, is **not complete**.
 - Consider Arad-Sibiu-Ard... (Infinite loop)
- Depth-first tree search can be modified at no extra memory cost so that **it checks new states against those on the path from the root to the current node**; this avoids infinite loops in finite state spaces but does not avoid the proliferation of redundant paths.

Depth-First Search (DFS)

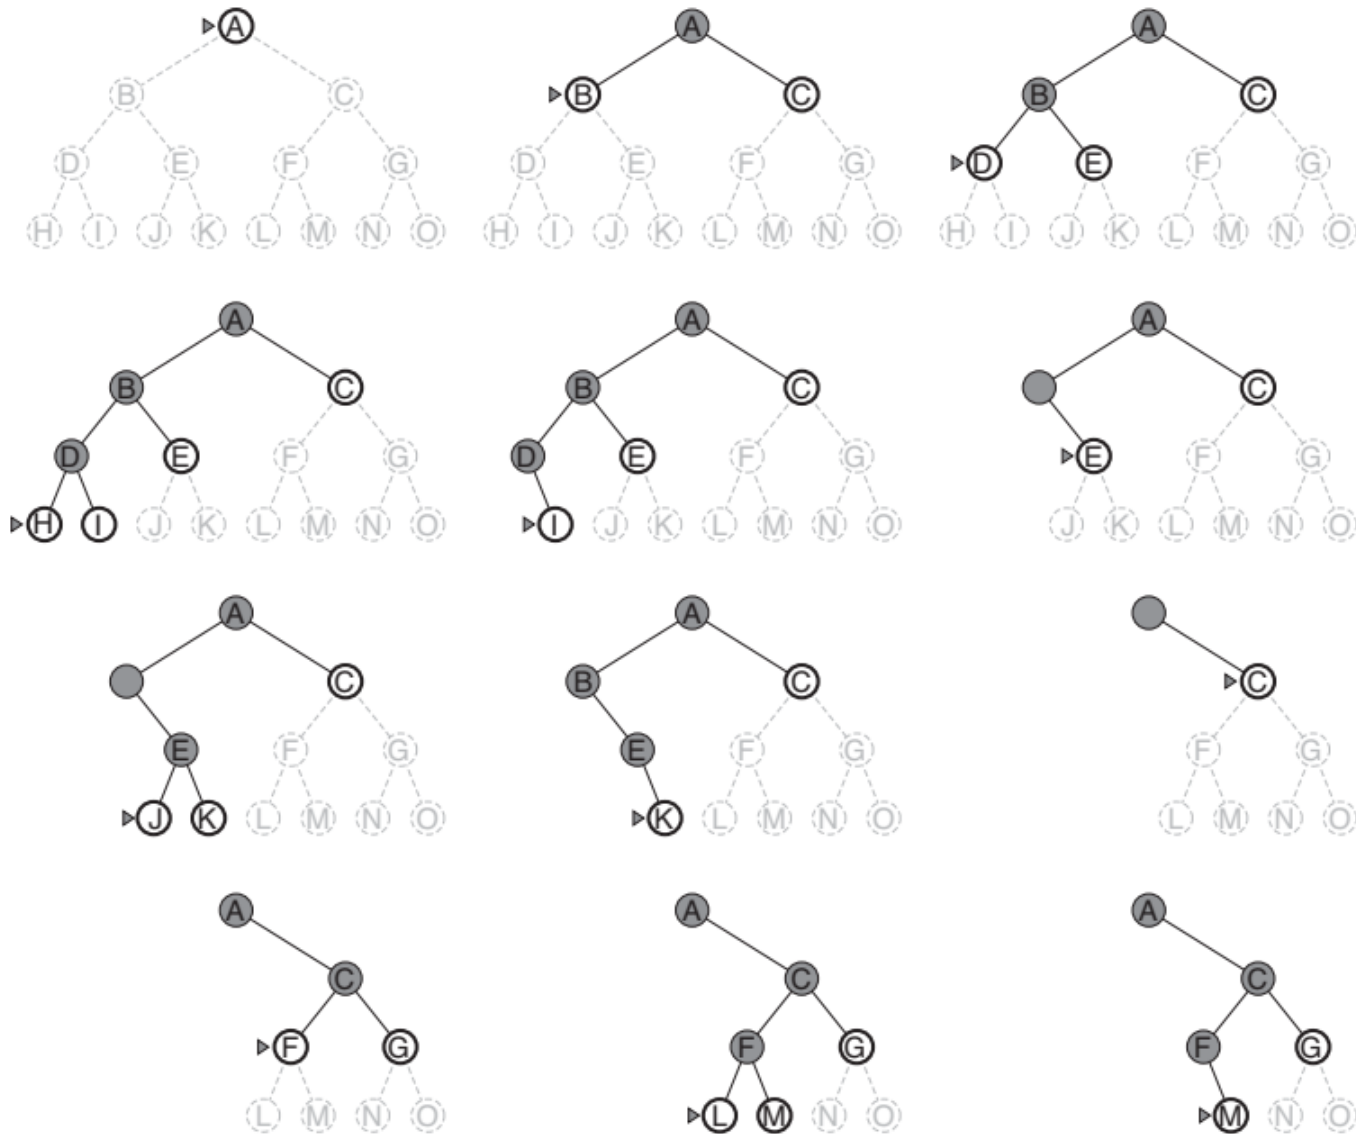
- DFS is not **optimal**.
 - In the example, it will go explore the entire left subtree even node C is a goal node.
 - If node J were also a goal node, then depth-first search would return it as a solution instead of C, which would be a better solution; hence, depth-first search is not optimal.
- Time Complexity: DFS tree may generate all of the $O(b^m)$ nodes in the tree.
 - m itself can be much larger than d (the depth of the shallowest solution) and is infinite if the tree is unbounded.
- Space Complexity: a depth-first tree search needs to **store only a single path** from the root to a leaf node, along with the remaining unexpanded sibling nodes for each node on the path.
 - Once a node has been expanded, it can be removed from memory as soon as all its descendants have been fully explored.
 - So, it requires $O(bm)$ nodes.

Depth-First Search (DFS)

```
function DEPTH-LIMITED-SEARCH(problem, limit) returns a solution, or failure/cutoff  
  return RECURSIVE-DLS(MAKE-NODE(problem.INITIAL-STATE), problem, limit)
```

```
function RECURSIVE-DLS(node, problem, limit) returns a solution, or failure/cutoff  
  if problem.GOAL-TEST(node.STATE) then return SOLUTION(node)  
  else if limit = 0 then return cutoff  
  else  
    cutoff_occurred?  $\leftarrow$  false  
    for each action in problem.ACTIONS(node.STATE) do  
      child  $\leftarrow$  CHILD-NODE(problem, node, action)  
      result  $\leftarrow$  RECURSIVE-DLS(child, problem, limit - 1)  
      if result = cutoff then cutoff_occurred?  $\leftarrow$  true  
      else if result  $\neq$  failure then return result  
  if cutoff_occurred? then return cutoff else return failure
```


Depth-First Search (DFS)



Depth-first search on a binary tree. The unexplored region is shown in light gray. Explored nodes with no descendants in the frontier are removed from memory. Nodes at depth 3 have no successors and M is the only goal node.

Depth-limited search

- The embarrassing failure of depth-first search in infinite state spaces can be **alleviated** by supplying depth-first search with a **predetermined depth limit l** .
 - In *depth-limited search*, nodes at depth l are treated as if they have no successors.
- The **depth limit** solves the infinite-path problem.
- It also introduces an additional source of incompleteness if we choose $l < d$, that is, the shallowest goal is beyond the depth limit.
- Its time complexity is $O(b^l)$ and its space complexity is $O(bl)$.
- Depth-first search can be viewed as a special case of depth-limited search with $l = \infty$.
- Sometimes, depth limits can be based on **knowledge** of the problem.
 - For example, on the map of Romania there are 20 cities. Therefore, we know that if there is a solution, it must be of length 19 at the longest, so $l = 19$ is a possible choice.
 - Also, we would discover that any city can be reached from any other city in at **most 9 steps**. This number, known as the **diameter** of the state space, gives us a better depth limit, which leads to a more efficient depth-limited search.

Depth-limited search

```
function DEPTH-LIMITED-SEARCH(problem, limit) returns a solution, or failure/cutoff  
  return RECURSIVE-DLS(MAKE-NODE(problem.INITIAL-STATE), problem, limit)
```

```
function RECURSIVE-DLS(node, problem, limit) returns a solution, or failure/cutoff  
  if problem.GOAL-TEST(node.STATE) then return SOLUTION(node)  
  else if limit = 0 then return cutoff  
  else  
    cutoff_occurred?  $\leftarrow$  false  
    for each action in problem.ACTIONS(node.STATE) do  
      child  $\leftarrow$  CHILD-NODE(problem, node, action)  
      result  $\leftarrow$  RECURSIVE-DLS(child, problem, limit - 1)  
      if result = cutoff then cutoff_occurred?  $\leftarrow$  true  
      else if result  $\neq$  failure then return result  
  if cutoff_occurred? then return cutoff else return failure
```

Iterative Deepening Depth-First Search

- **Iterative deepening search** (or iterative deepening depth-first search) is a general strategy, often used in combination with **depth-first tree** search, that finds the best depth limit.
 - It does this by **gradually increasing the limit**—first 0, then 1, then 2, and so on—until a goal is found.
 - This will occur when the depth limit reaches d , the depth of the shallowest goal node.
- Iterative deepening combines the benefits of **depth-first** and **breadth-first** search.
- Space complexity: $O(bd)$.
- Time complexity: Same as BFS, i.e. $O(b^d)$.
- Completeness: like breadth-first search, it is **complete** when the branching factor is finite and **optimal** when the path cost is a nondecreasing function of the depth of the node.
- Iterative deepening search may seem wasteful because states are generated multiple times. It turns out this is not too costly.
 - The reason is that in a search tree with the same (or nearly the same) branching factor at each level, most of the nodes are in the bottom level, **so it does not matter much that the upper levels are generated multiple times**.

Iterative Deepening Depth-First Search

```
function ITERATIVE-DEEPENING-SEARCH(problem) returns a solution, or failure
  for depth = 0 to  $\infty$  do
    result  $\leftarrow$  DEPTH-LIMITED-SEARCH(problem, depth)
    if result  $\neq$  cutoff then return result
```

Iterative Deepening Depth-First Search

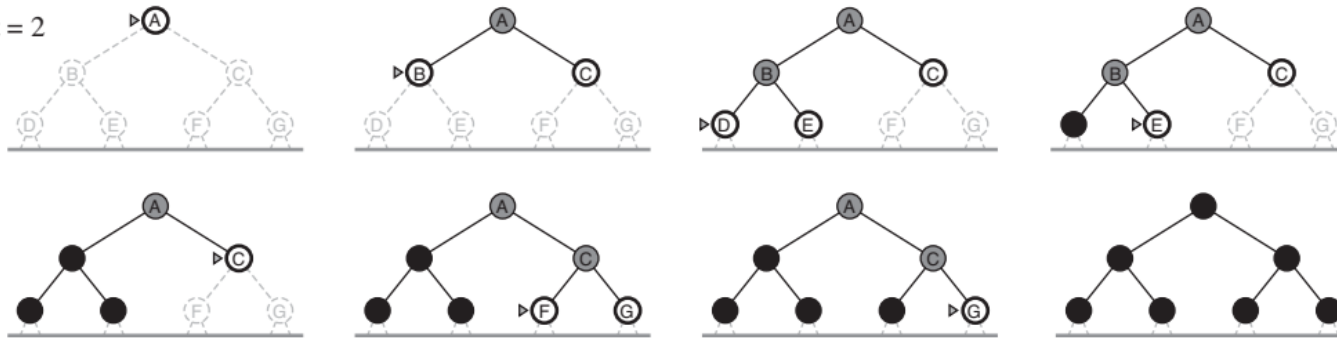
Limit = 0



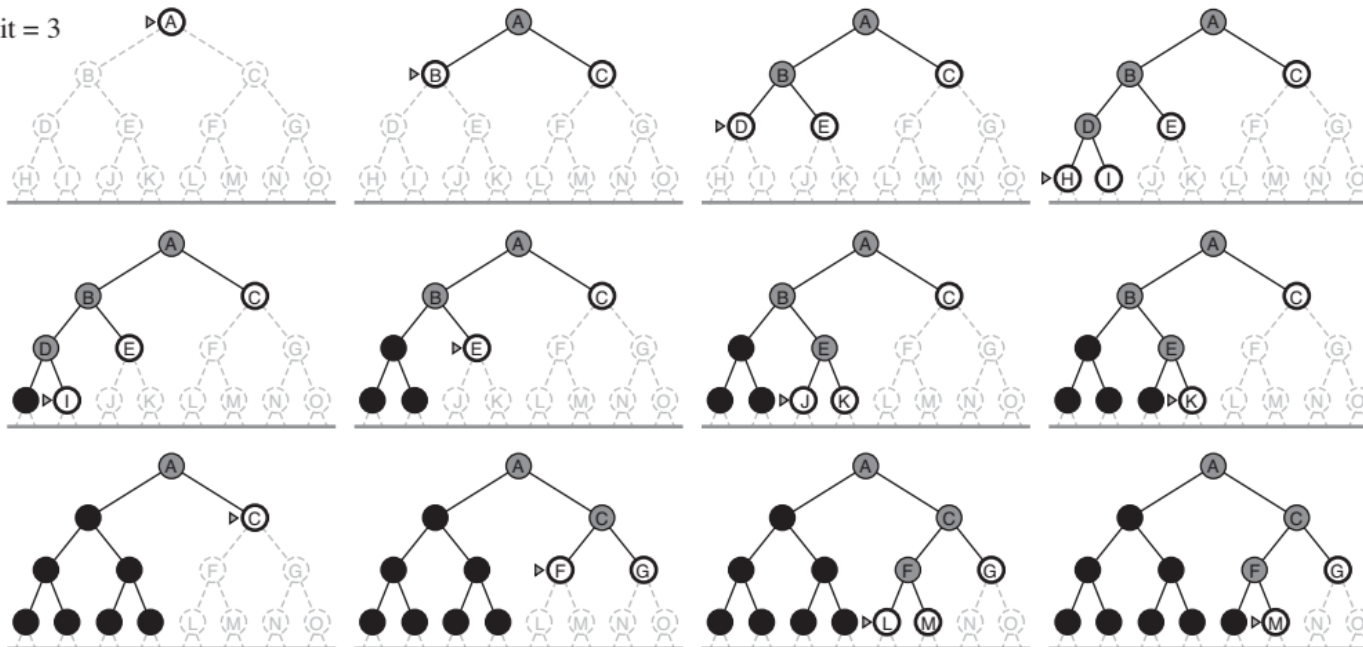
Limit = 1



Limit = 2



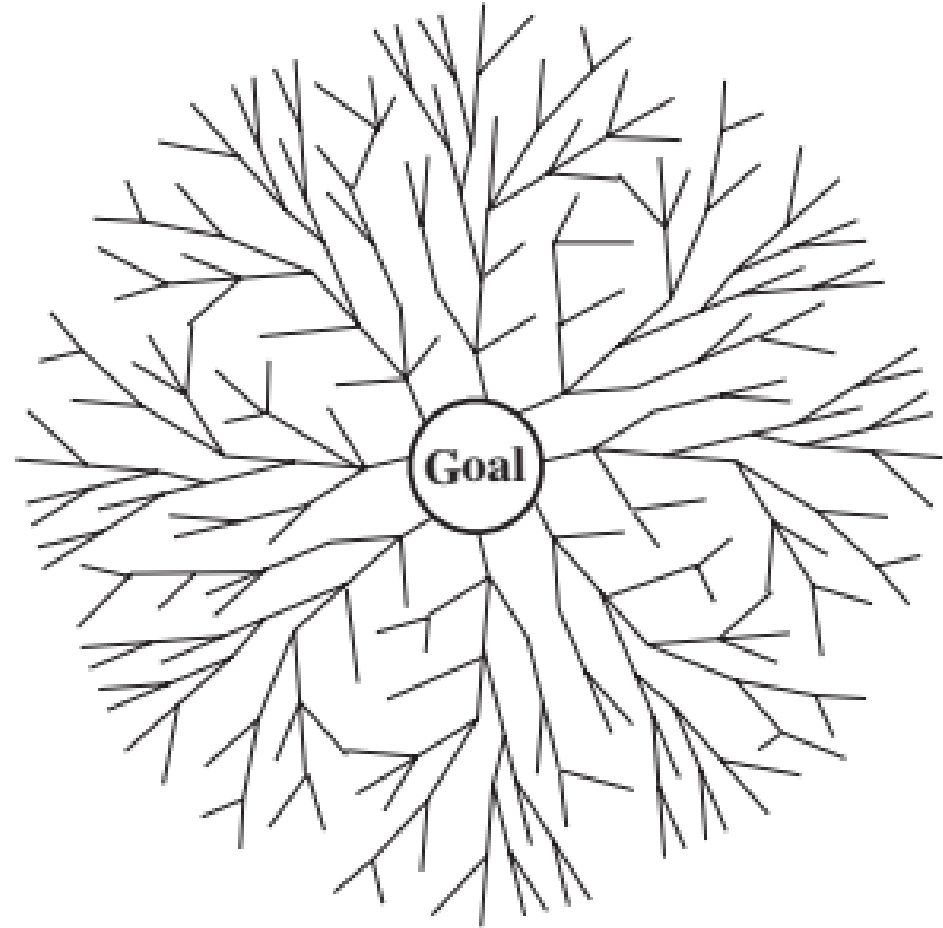
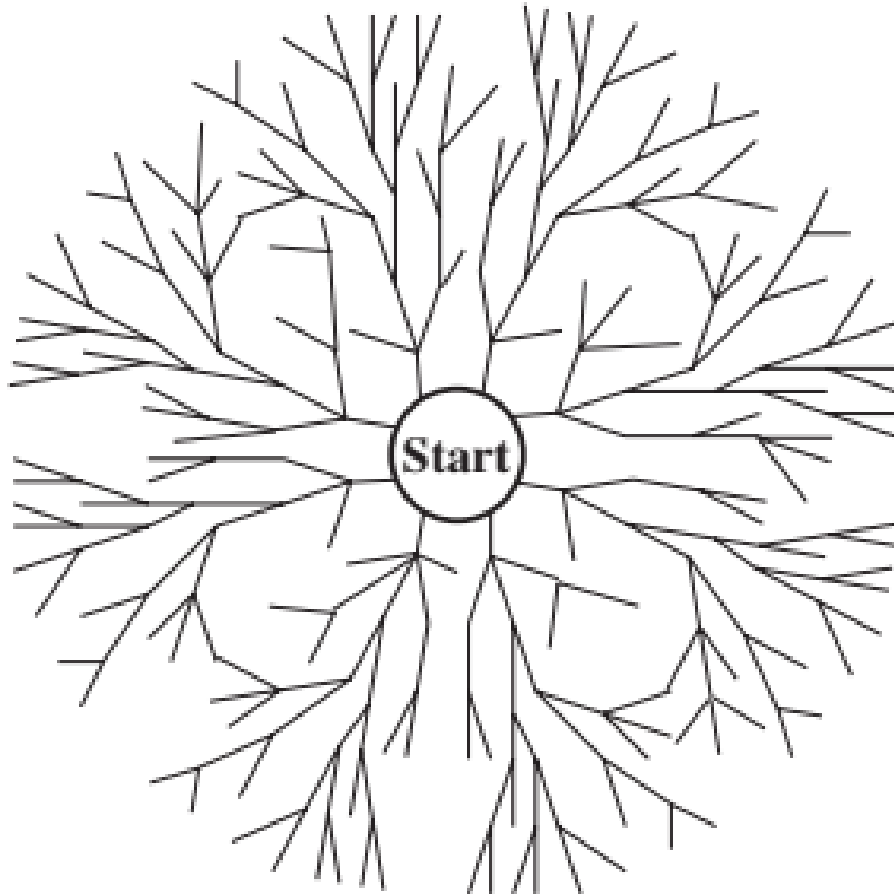
Limit = 3



Bidirectional search

- The idea behind **bidirectional search** is to **run two simultaneous searches**—one forward from the **initial state** and the other backward from the **goal**—hoping that the two searches **meet** in the middle.
- Time and complexity: $b^{d/2} + b^{d/2} = O(b^{d/2})$
- **Bidirectional search** is implemented by replacing the goal test with **a check to see whether the frontiers of the two searches intersect**; if they do, a solution has been found.
- Consider the question of what we mean by “the goal” in searching “backward from the goal.” For the 8-puzzle and for finding a route in Romania, there is just **one goal state**, so the backward search is very much like the forward search.
- If there are several explicitly listed goal states—for example, the two dirt-free goal states—then we can construct a new dummy goal state whose immediate predecessors are all the actual goal states.
 - But if the goal is an **abstract description**, such as the goal that “no queen attacks another queen” in the n-queens problem, then bidirectional search is **difficult** to use.

Bidirectional search



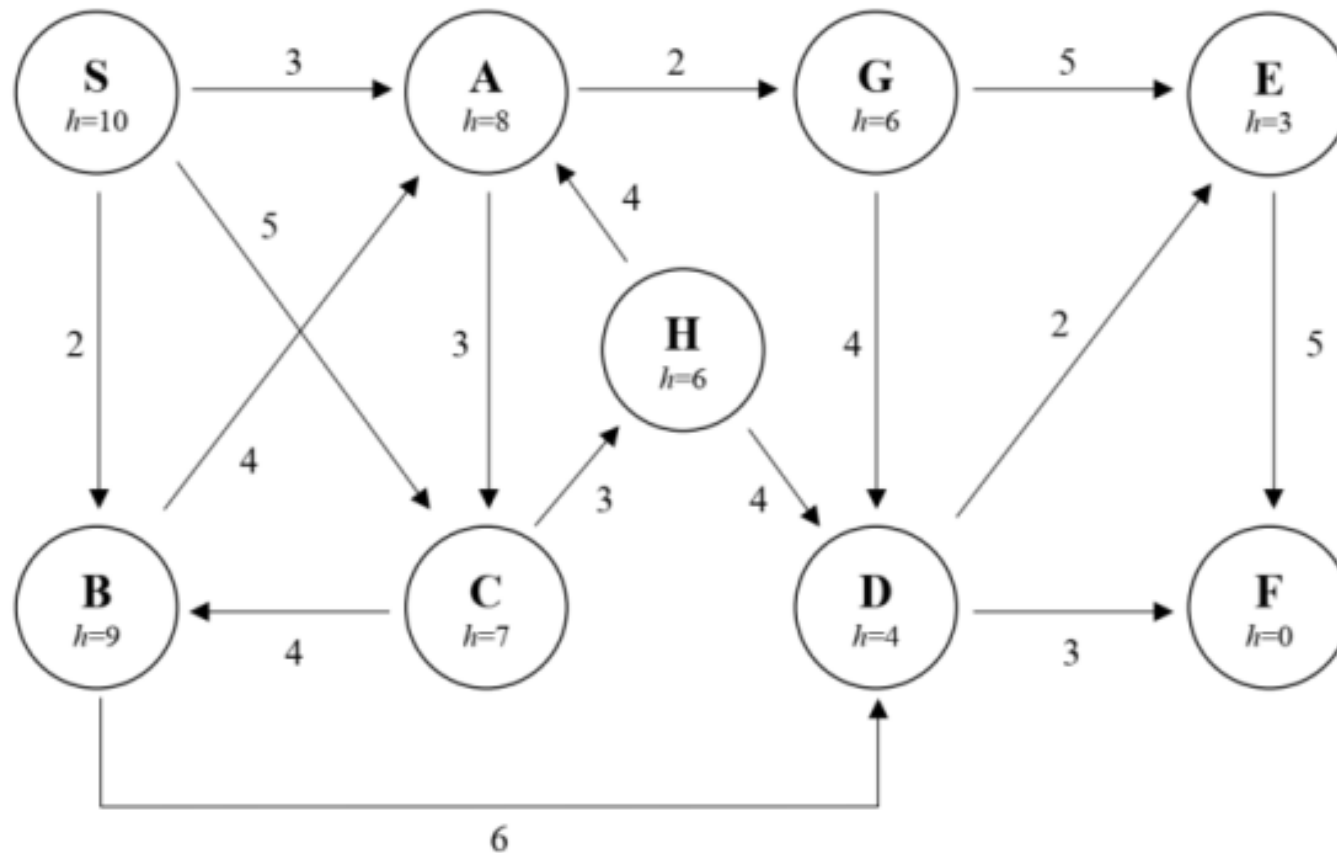
Comparison

Criterion	Breadth-First	Uniform-Cost	Depth-First	Depth-Limited	Iterative Deepening	Bidirectional (if applicable)
Complete?	Yes ^a	Yes ^{a,b}	No	No	Yes ^a	Yes ^{a,d}
Time	$O(b^d)$	$O(b^{1+\lceil C^*/\epsilon \rceil})$	$O(b^m)$	$O(b^\ell)$	$O(b^d)$	$O(b^{d/2})$
Space	$O(b^d)$	$O(b^{1+\lceil C^*/\epsilon \rceil})$	$O(bm)$	$O(b\ell)$	$O(bd)$	$O(b^{d/2})$
Optimal?	Yes ^c	Yes	No	No	Yes ^c	Yes ^{c,d}

Evaluation of tree-search strategies. b is the branching factor; d is the depth of the shallowest solution; m is the maximum depth of the search tree; ℓ is the depth limit. Superscript caveats are as follows: ^a complete if b is finite; ^b complete if step costs $\geq \epsilon$ for positive ϵ ; ^c optimal if step costs are all identical; ^d if both directions use breadth-first search.

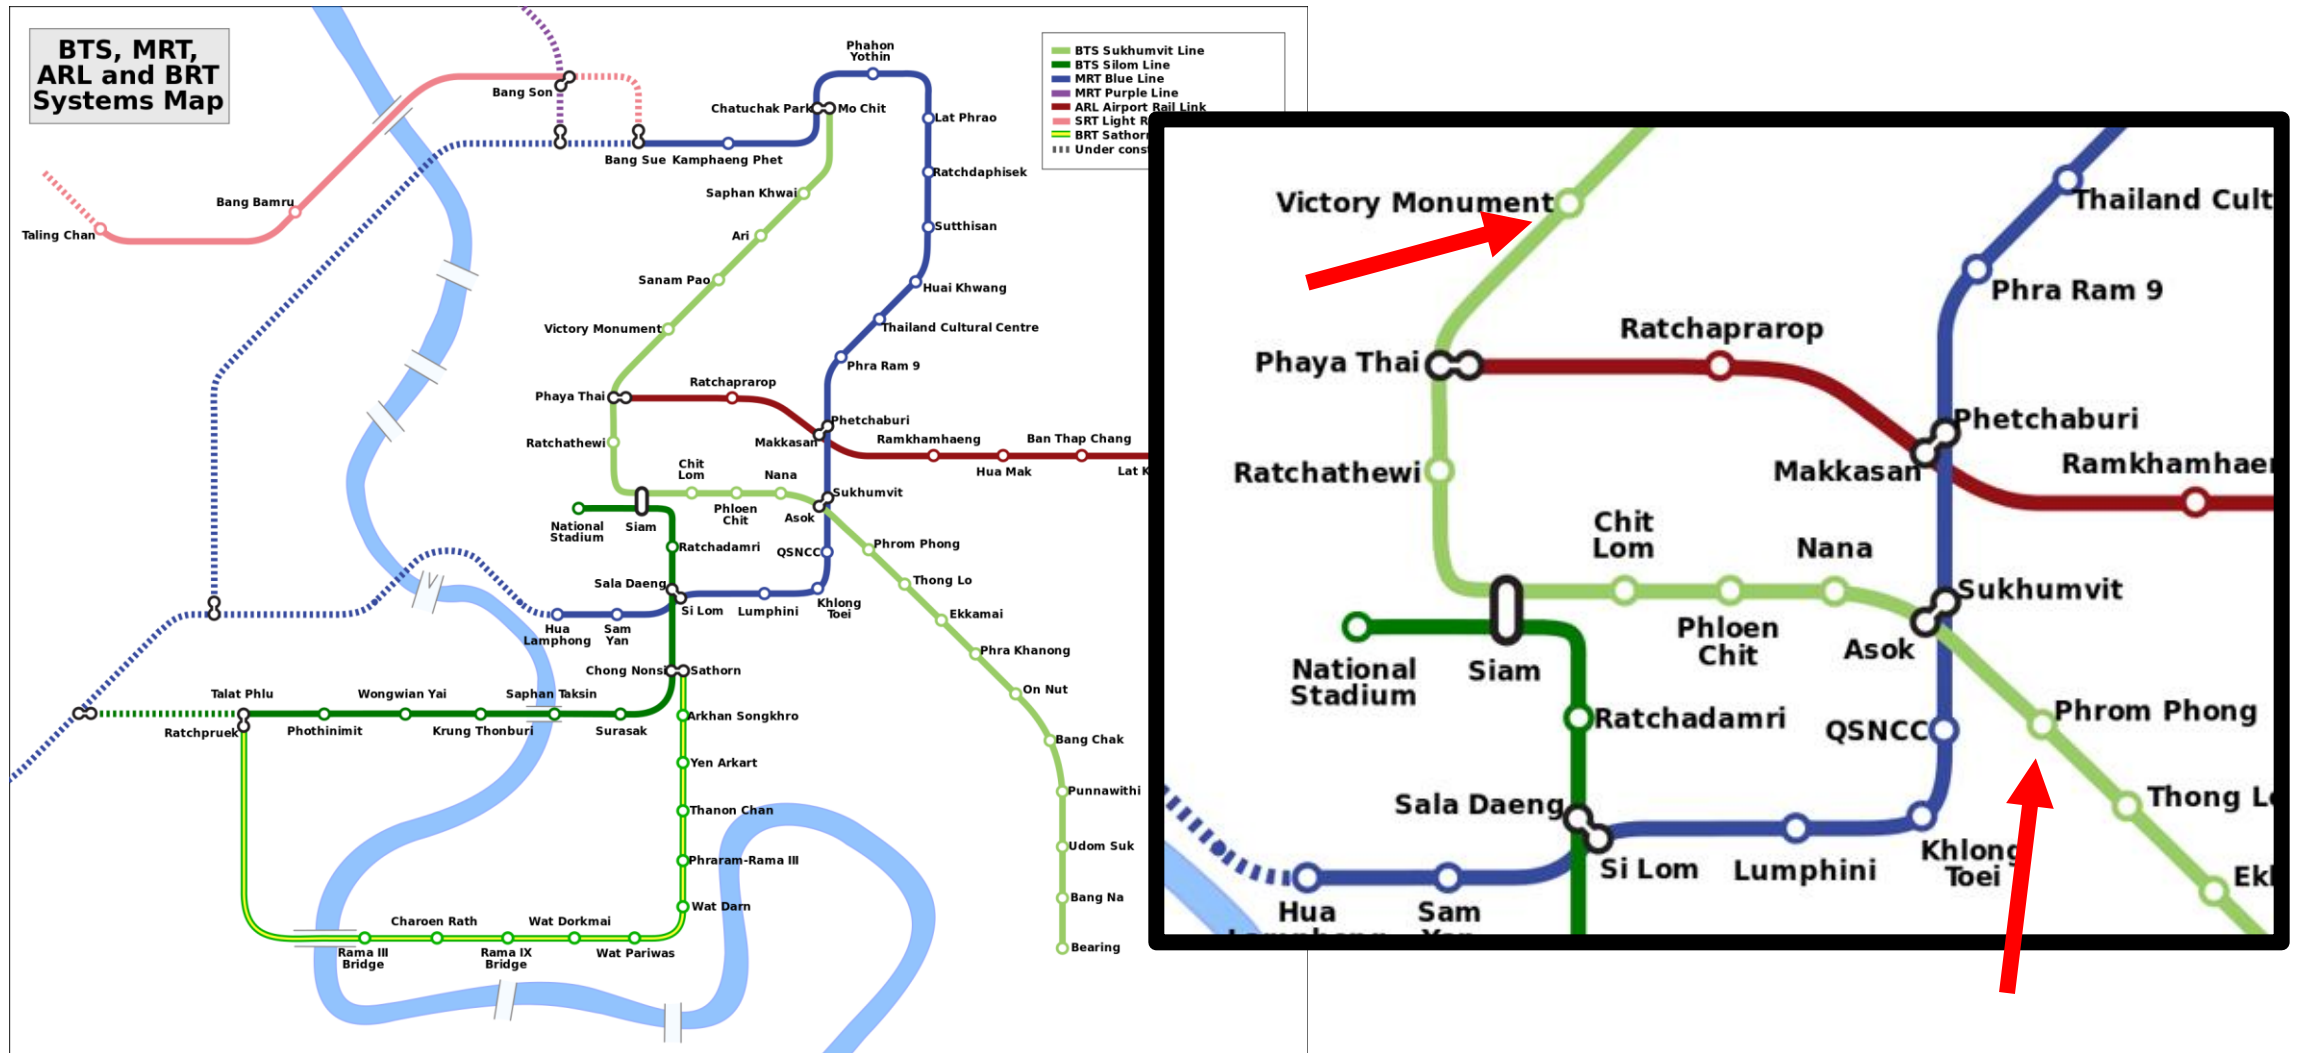
Exercise

Problem: ให้นักเรียนเขียนลำดับการ visit nodes ตามความเข้าใจ เมื่อทำเสร็จแล้ว
ให้เช็คคำตอบกับผลลัพธ์จากการรันโปรแกรมของตนเอง



Single-Source Shortest Paths

Single-Source Shortest Paths



Shortest-paths problem

Note: Edge weights can represent metrics other than distances, such as time, cost, penalties, loss, or any other quantity that accumulates linearly along a path and that we would want to minimize

- In a shortest-paths problem, we are given a weighted, directed graph $G(V, E)$, with weight function $w: E \rightarrow R$ mapping edges to real-valued weights.
- The weight $w(p)$ of path $p = \langle v_0, v_1, \dots, v_k \rangle$ is the sum of the weights of its constituent edges: $w(p) = \sum_{i=1}^k w(v_{i-1}, v_i)$
- We define the shortest-path weight $\delta(u, v)$ from u to v by

$$\delta(u, v) = \begin{cases} \min\{w(p) : u \xrightarrow{p} v\} & \text{if there is a path from } u \text{ to } v, \\ \infty & \text{otherwise.} \end{cases}$$

- A **shortest path** from vertex u to vertex v is then defined as any path p with weight $w(p) = \delta(u, v)$.

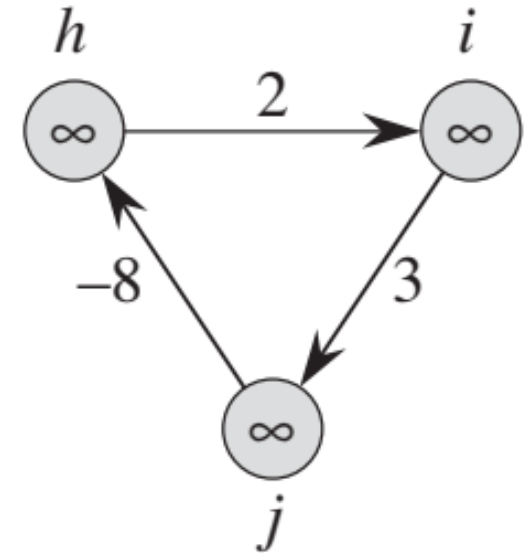
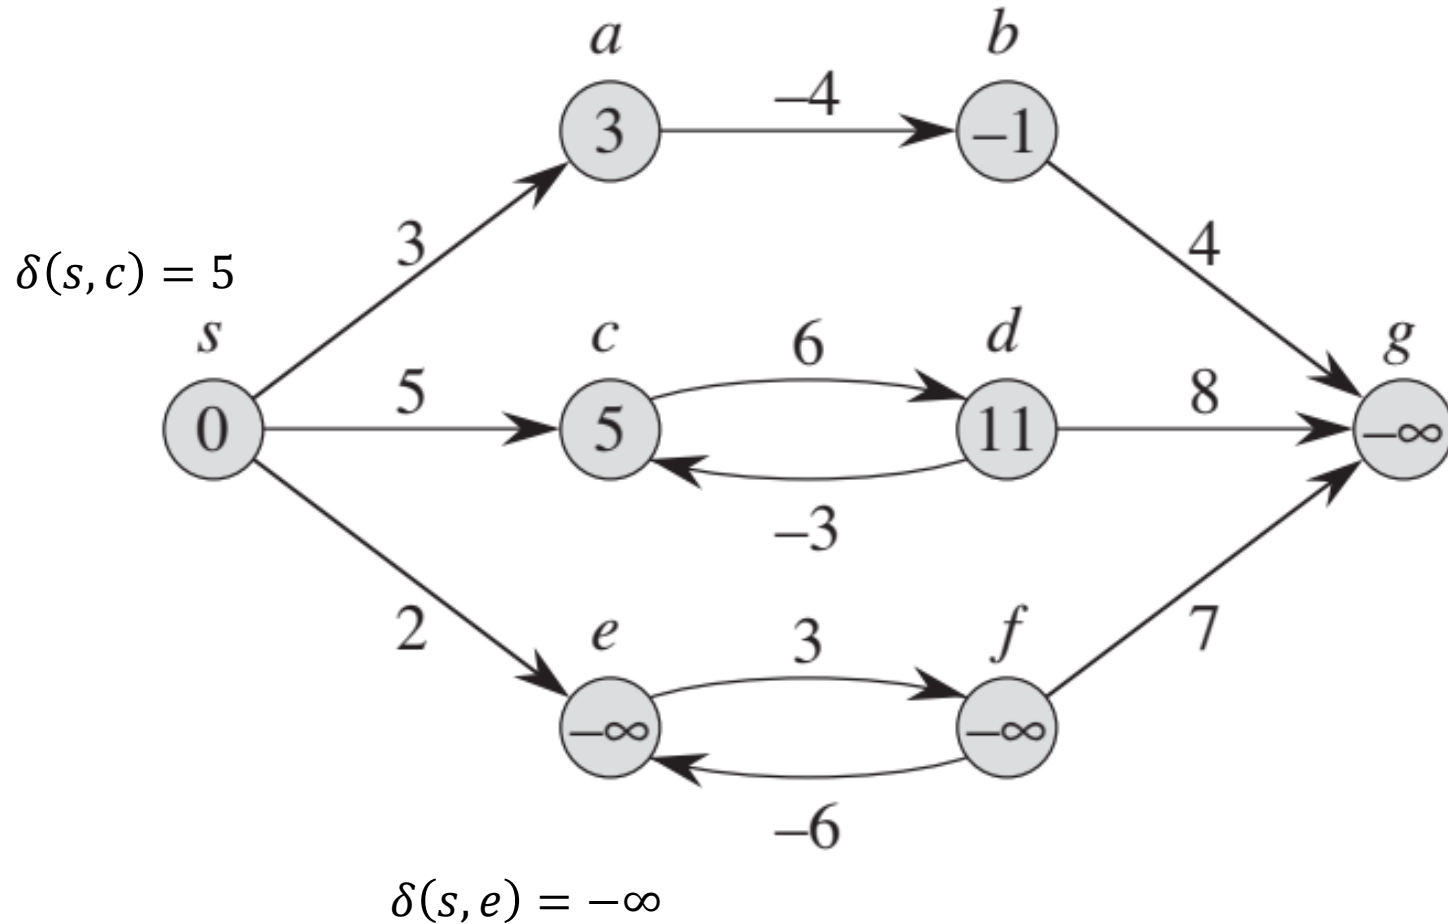
Optimal substructure of a shortest path

- Shortest-paths algorithms typically rely on the property that a shortest path between two vertices contains other shortest paths within it.
- Recall that optimal substructure is one of the key indicators that **dynamic programming** and the **greedy method** might apply.
- **Dijkstra's algorithm** is a greedy algorithm.

Negative-weight edges

- Some instances of the single-source shortest-paths problem may include edges whose weights are **negative**.
- If the graph $G(V, E)$ contains **no negative weight cycles** reachable from the source s , then for all $v \in V$, the shortest-path weight $\delta(s, v)$ remains well defined, even if it has a negative value.
- If the graph contains a **negative-weight cycle** reachable from s , however, shortest-path weights are not well defined.

Negative-weight edges



Cycles

- Can a shortest path contain a cycle? As we have just seen, it cannot contain a **negative-weight cycle**.
- How about **positive-weight cycle**?
 - Nor can it contain a positive-weight cycle, since removing the cycle from the path produces a path with the same source and destination vertices and a lower path weight.
- Therefore, without loss of generality we can assume that when we are finding shortest paths, they have no cycles, i.e., they are simple paths.

Relaxation

- For each vertex $v \in V$, we maintain an attribute $v.d$, which is an upper bound on the weight of a shortest path from source s to v .
 - We call $v.d$ a **shortest-path estimate**.
- We initialize the **shortest-path estimates** and **predecessors** by the following $\Theta(V)$ -time procedure:

INITIALIZE-SINGLE-SOURCE(G, s)

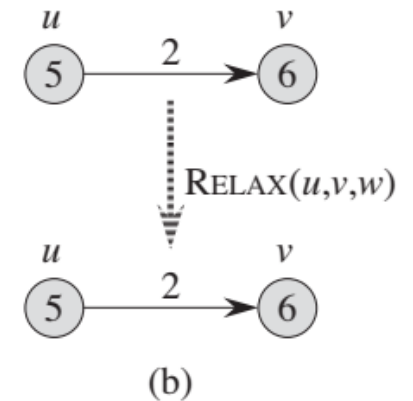
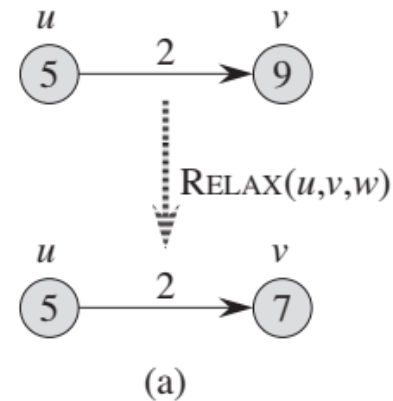
```
1  for each vertex  $v \in G.V$ 
2       $v.d = \infty$ 
3       $v.\pi = \text{NIL}$ 
4   $s.d = 0$ 
```

Relaxation

- The process of relaxing an edge (u, v) consists of testing whether we can improve the shortest path to v found so far by going through u and, if so, updating $v.d$ and $v.\pi$.
- The following code performs a relaxation step on edge (u, v) in $O(1)$ time:

RELAX(u, v, w)

```
1  if  $v.d > u.d + w(u, v)$   
2       $v.d = u.d + w(u, v)$   
3       $v.\pi = u$ 
```



The Bellman-Ford algorithm

The Bellman-Ford algorithm

- The **Bellman-Ford algorithm** solves the single-source shortest-paths problem in the general case in which edge weights may be **negative**.
- Given a weighted, directed graph $G = (V, E)$ with source s and weight function $w: E \rightarrow R$, the Bellman-Ford algorithm returns a Boolean value indicating whether or not there is a **negative-weight cycle** that is reachable from the **source**.
 - If there is such a cycle, the algorithm indicates that no solution exists.
 - If there is no such cycle, the algorithm produces the shortest paths and their weights.
- The Bellman-Ford algorithm relaxes each edge $|V| - 1$ times.
- The algorithm relaxes edges, progressively decreasing an estimate $v.d$ on the weight of a shortest path from the source s to each vertex $\delta(s, v)$ until it achieves the actual shortest-path weight

The Bellman-Ford algorithm

BELLMAN-FORD(G, w, s)

```
1  INITIALIZE-SINGLE-SOURCE( $G, s$ )
2  for  $i = 1$  to  $|G.V| - 1$ 
3      for each edge  $(u, v) \in G.E$ 
4          RELAX( $u, v, w$ )
5  for each edge  $(u, v) \in G.E$ 
6      if  $v.d > u.d + w(u, v)$ 
7          return FALSE
8  return TRUE
```

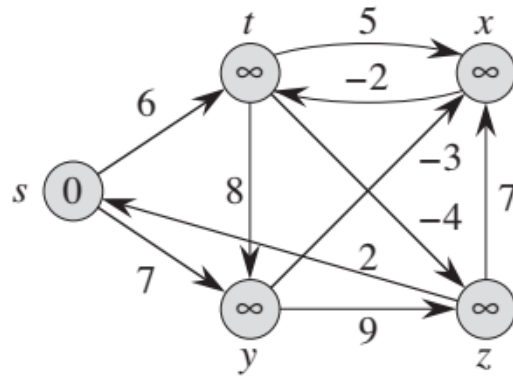
INITIALIZE-SINGLE-SOURCE(G, s)

```
1  for each vertex  $v \in G.V$ 
2       $v.d = \infty$ 
3       $v.\pi = \text{NIL}$ 
4   $s.d = 0$ 
```

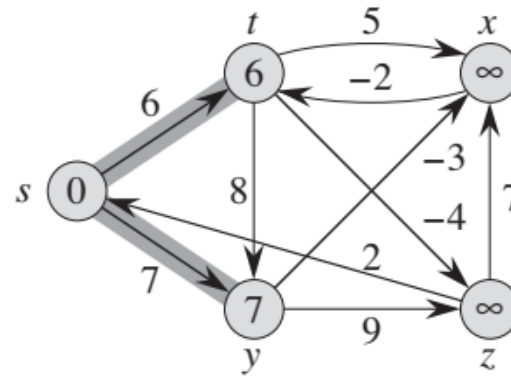
RELAX(u, v, w)

```
1  if  $v.d > u.d + w(u, v)$ 
2       $v.d = u.d + w(u, v)$ 
3       $v.\pi = u$ 
```

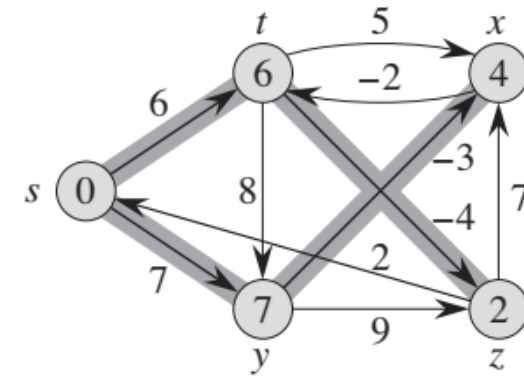
The Bellman-Ford algorithm



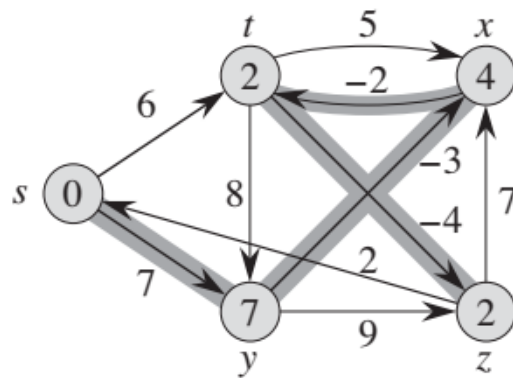
(a)



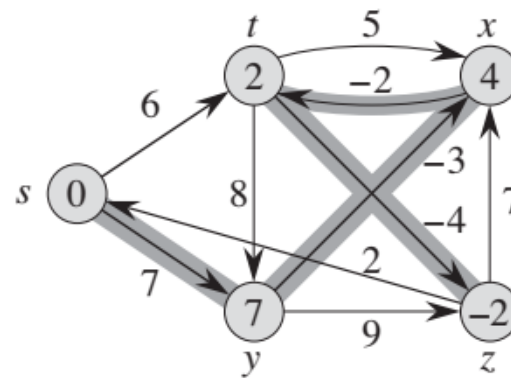
(b)



(c)



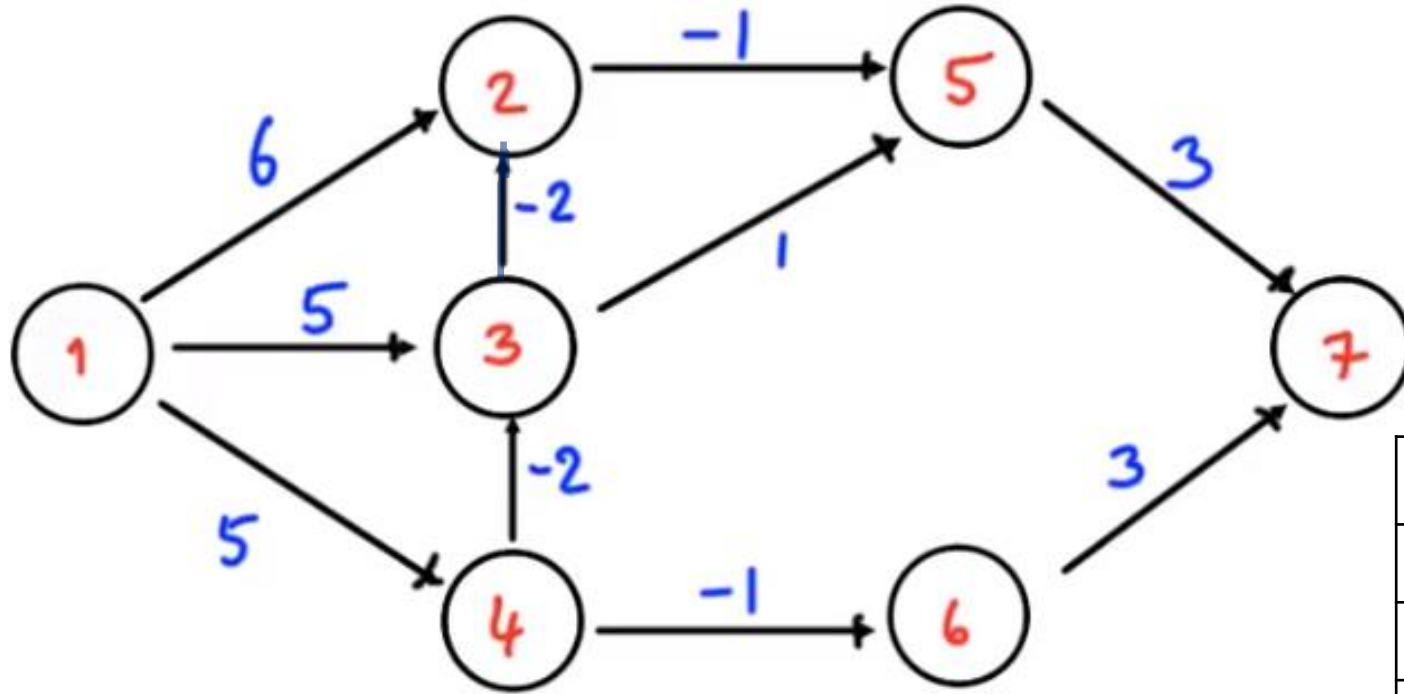
(d)



(e)

Quiz#08

Source: node 1

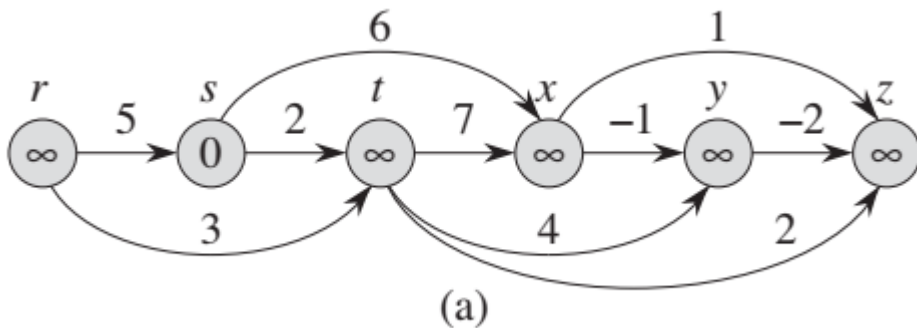


Complexity Analysis

- The Bellman-Ford algorithm runs in time $O(VE)$, since the initialization takes $\Theta(V)$ time, each of the $|V| - 1$ passes over the edges takes $O(E)$ time.

Single-source shortest paths in directed acyclic graphs

- By relaxing the edges of a weighted **dag** (directed acyclic graph) $G = (V, E)$ according to a topological sort of its vertices, we can compute shortest paths from a single source in $\Theta(V + E)$ time.
- Shortest paths are always well defined in a dag, since even if there are negative-weight edges, **no negative-weight cycles can exist**.
- The algorithm starts by **topologically sorting** the dag to impose a linear ordering on the vertices.



DAG-SHORTEST-PATHS(G, w, s)

```
1  topologically sort the vertices of  $G$ 
2  INITIALIZE-SINGLE-SOURCE( $G, s$ )
3  for each vertex  $u$ , taken in topologically sorted order
4      for each vertex  $v \in G.Adj[u]$ 
5          RELAX( $u, v, w$ )
```

Dijkstra's algorithm

Dijkstra's algorithm

- Dijkstra's algorithm solves the single-source shortest-paths problem on a weighted, directed graph $G = (V, E)$ for the case in which all edge weights are **nonnegative**.
- As we shall see, with a good implementation, the running time of Dijkstra's algorithm is lower than that of the Bellman-Ford algorithm.
- Dijkstra's algorithm maintains a set S of vertices whose final shortest-path weights from the source s have already been determined.
- The algorithm repeatedly selects the vertex $u \in V - S$ with the **minimum shortest-path estimate**, adds u to S , and relaxes all edges leaving u .

Dijkstra's algorithm

DIJKSTRA(G, w, s)

```
1  INITIALIZE-SINGLE-SOURCE( $G, s$ )
2   $S = \emptyset$ 
3   $Q = G.V$ 
4  while  $Q \neq \emptyset$ 
5       $u = \text{EXTRACT-MIN}(Q)$ 
6       $S = S \cup \{u\}$ 
7      for each vertex  $v \in G.Adj[u]$ 
8          RELAX( $u, v, w$ )
```

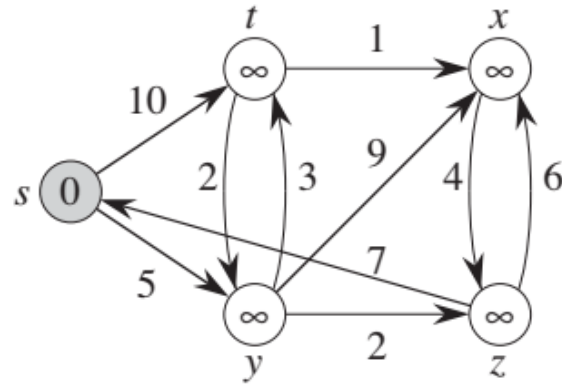
INITIALIZE-SINGLE-SOURCE(G, s)

```
1  for each vertex  $v \in G.V$ 
2       $v.d = \infty$ 
3       $v.\pi = \text{NIL}$ 
4   $s.d = 0$ 
```

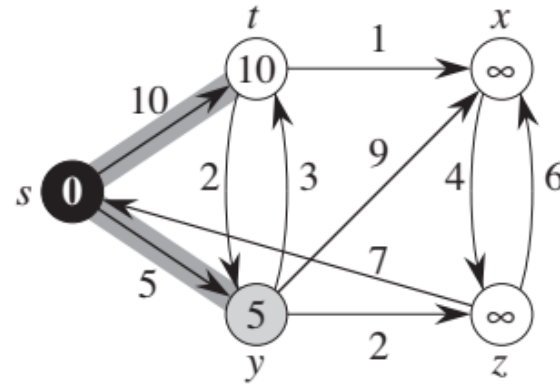
RELAX(u, v, w)

```
1  if  $v.d > u.d + w(u, v)$ 
2       $v.d = u.d + w(u, v)$ 
3       $v.\pi = u$ 
```

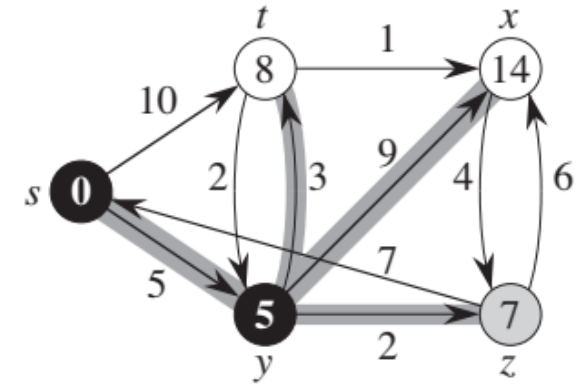
Dijkstra's algorithm



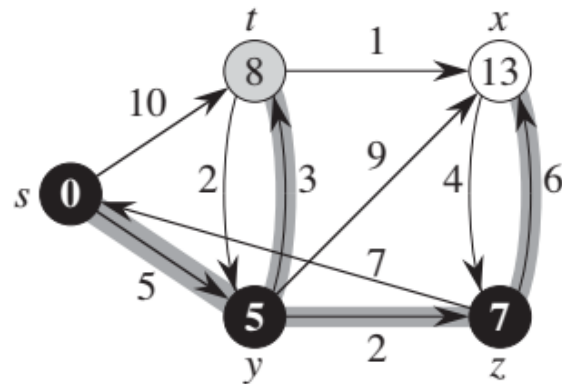
(a)



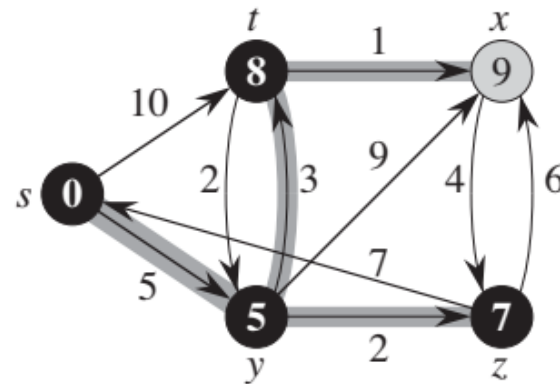
(b)



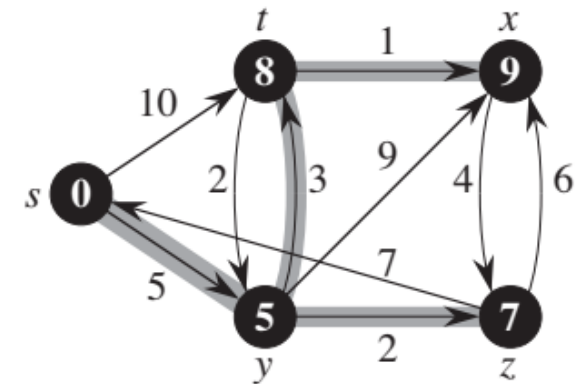
(c)



(d)

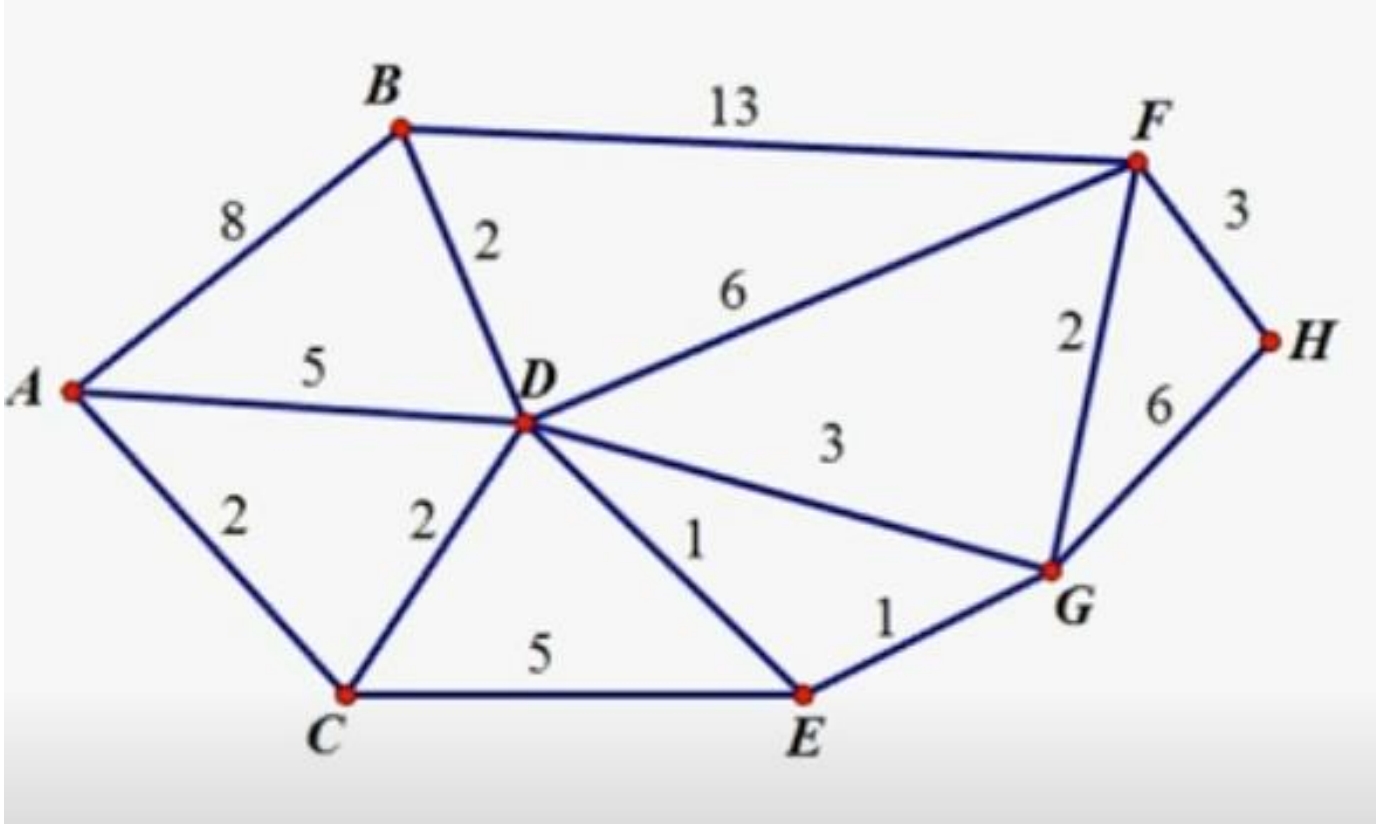


(e)



(f)

Exercise



โครงสร้างข้อมูลแบบเซตไม่มี ส่วนร่วม **Disjoint Sets**

Disjoint-set

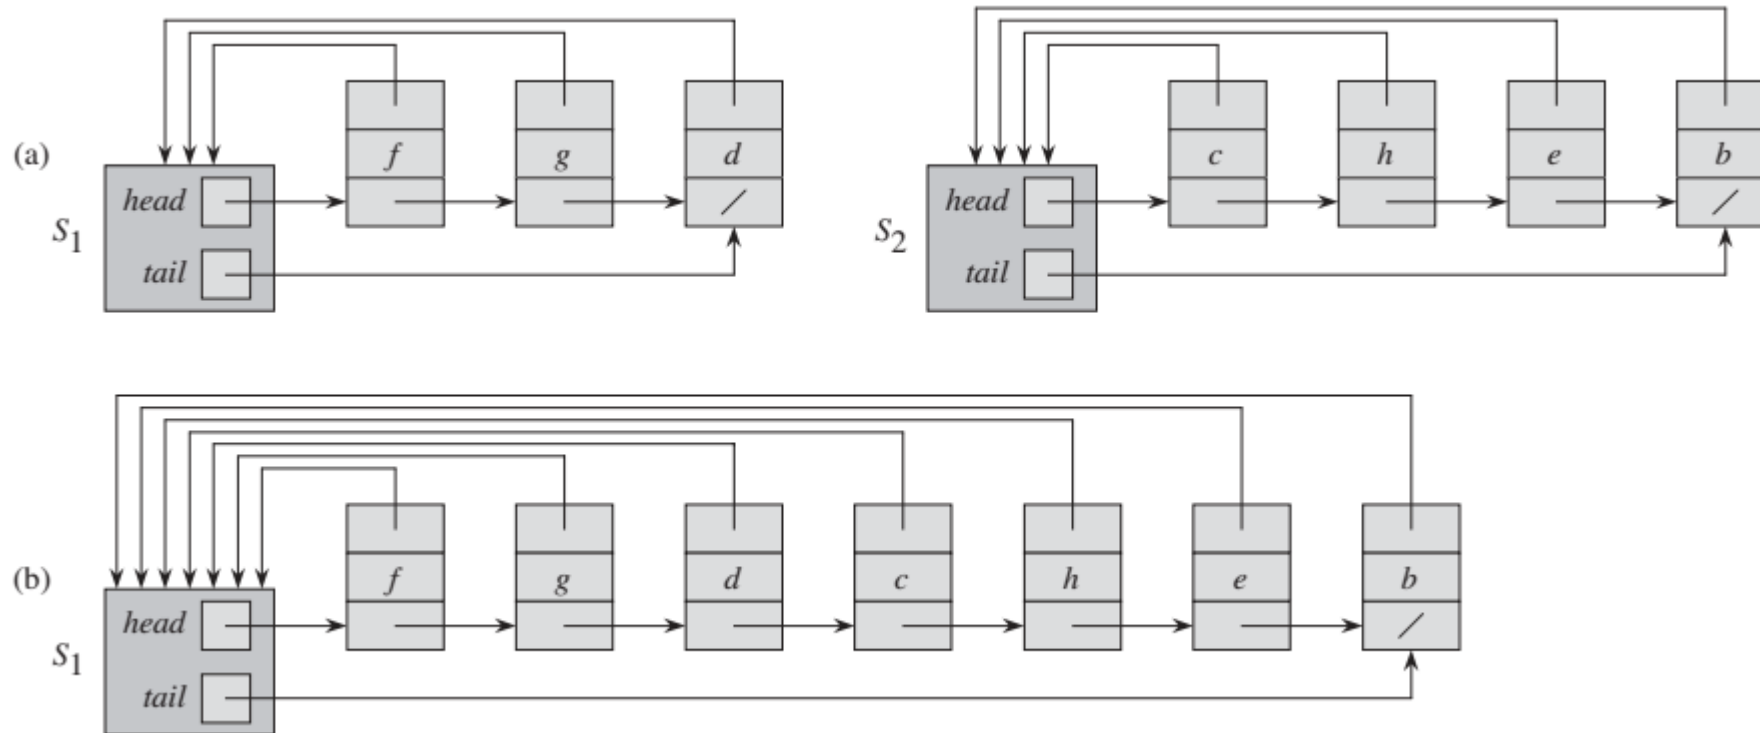
- A disjoint-set data structure maintains a collection $\mathcal{S} = \{S_1, \dots, S_k\}$ of disjoint **dynamic** sets.
- We identify each set by a **representative**, which is some member of the set.

Disjoint-set Operations

- As in the other dynamic-set implementations we have studied, we represent each element of a set by an object. Letting x denote an object, we wish to support the following operations:
 - **MAKE-SET(x)** creates a new set whose only member (and thus representative) is x . Since the sets are disjoint, we require that x not already be in some other set.
 - **UNION(x, y)** unites the dynamic sets that contain x and y , say S_x and S_y , into a new set that is the union of these two sets.
 - **FIND-SET(x)** returns a pointer to the representative of the (unique) set containing x .

Implementation

- Linked List representation



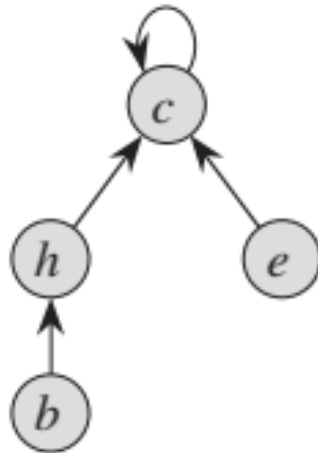
Implementation

- Forest Representation.

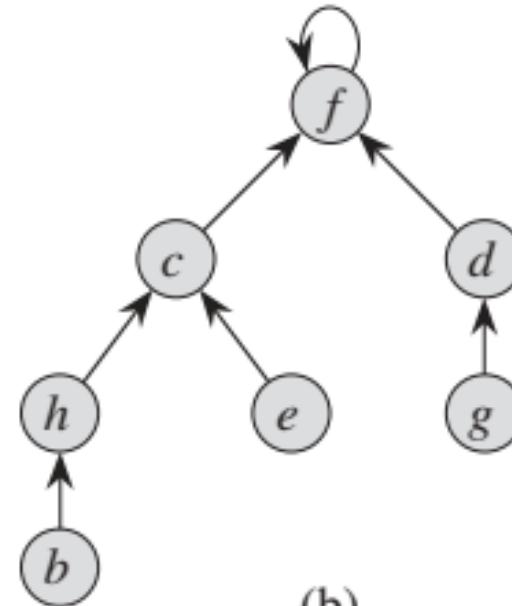
MAKE-SET(x) create a new tree with one node in the forest.

UNION(x,y) following parent pointers until we find the root of the tree.

FIND-SET(x) causes the root of one tree to point to the root of the other.



(a)

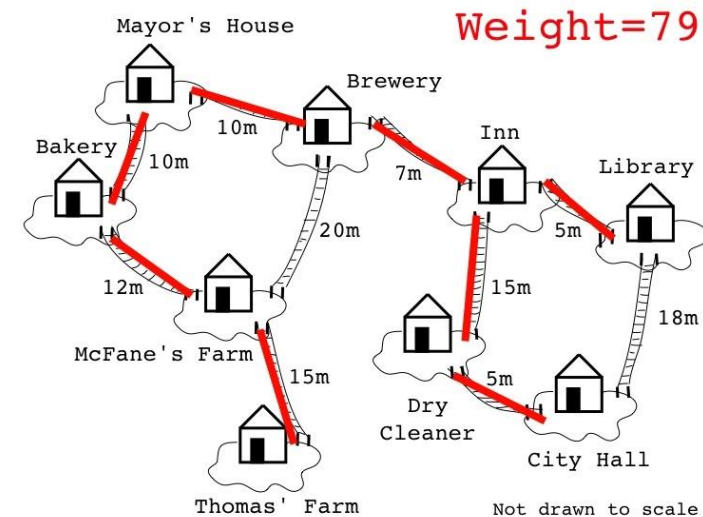


(b)

Minimum Spanning Trees

Minimum Spanning Trees

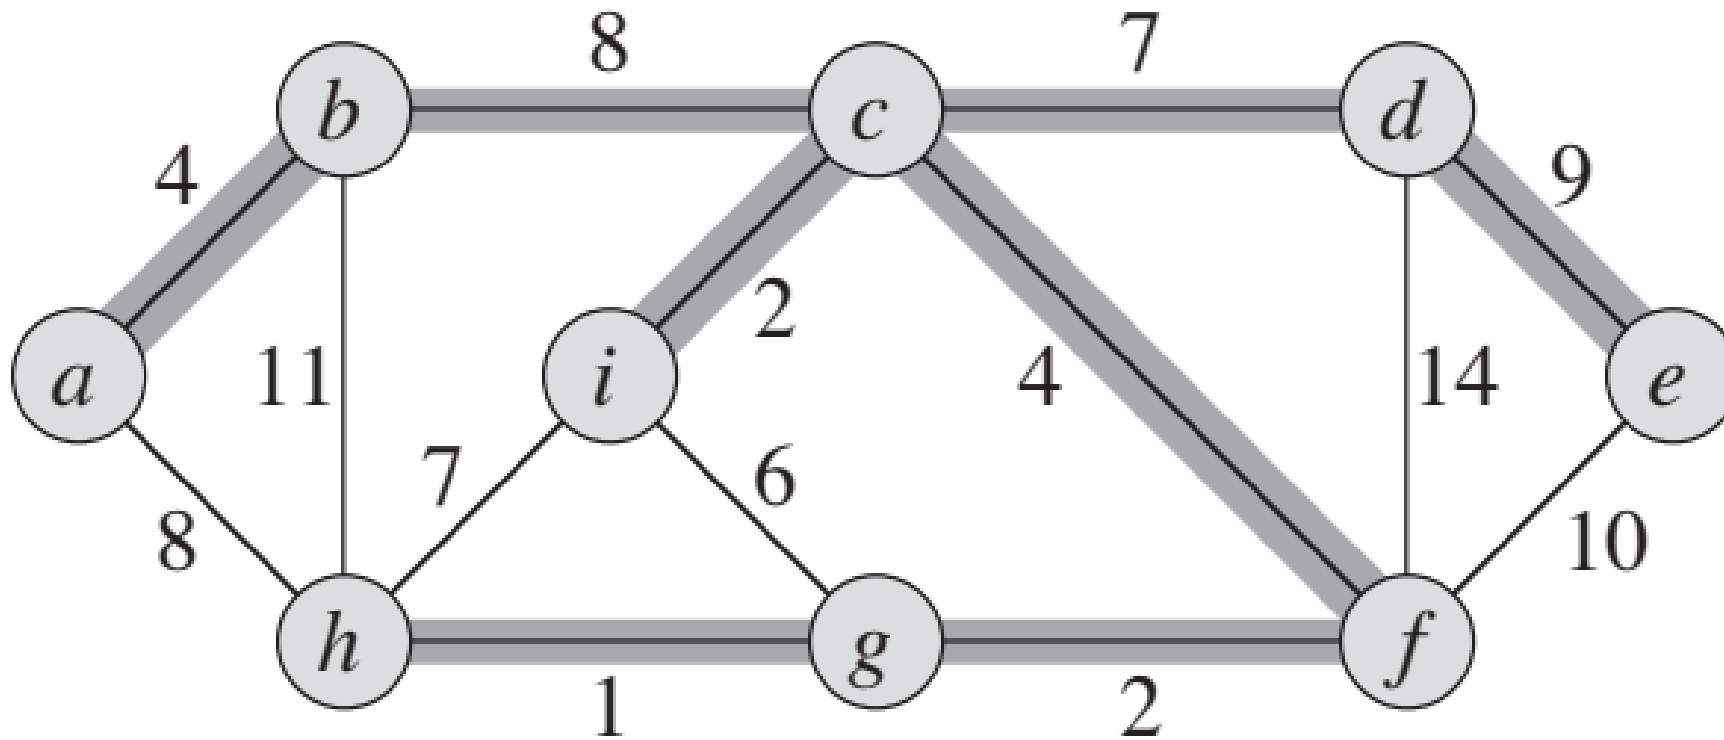
- Electronic circuit designs often need to make the pins of several components electrically equivalent by wiring them together.
- To interconnect a set of n pins, we can use an arrangement of $n - 1$ wires, each connecting two pins.
- Of all such arrangements, the one that uses the **least amount of wire** is usually the most desirable.



Minimum Spanning Trees

- We can model this wiring problem with a connected, undirected graph $G = (V, E)$, where V is the **set of pins**. E is the set of **possible interconnections between pairs of pins**, and for each edge $(u, v) \in E$, we have a weight $w(u, v)$ specifying the cost (amount of wire needed) to connect u and v .
- In other words, We then wish to find an acyclic subset $T \subseteq E$, that connects **all of the vertices** and whose total weight $w(T) = \sum_{(u,v) \in T} w(u, v)$
- Since T is **acyclic** and **connects** all of the vertices, it must form a tree, which we call a spanning tree since it “spans” the graph G .
- We call the problem of determining the tree T the **minimum-spanning-tree** problem.

Minimum Spanning Trees



Minimum Spanning Trees

- In this lecture, we shall examine two algorithms for solving the minimum spanning-tree problem: **Kruskal**'s algorithm and **Prim**'s algorithm.
- We can easily make each of them run in time $O(E \lg V)$ using ordinary binary heaps.

Growing a minimum spanning tree

Growing a minimum spanning tree

- Assume that we have a *connected, undirected* graph $G = (V, E)$ with a weight function $w: E \rightarrow R$, and we wish to find a **minimum spanning tree** for G .
- Prior to each iteration, A is a subset of some minimum spanning tree.
- At each step, we determine an edge (u, v) that we can add to A without violating this invariant, so $A \cup \{(u, v)\}$ is also a subset of a minimum spanning tree.
 - We call such an edge a safe edge for A , since we can add it safely to A while maintaining the invariant.

GENERIC-MST(G, w)

```
1   $A = \emptyset$ 
2  while  $A$  does not form a spanning tree
3      find an edge  $(u, v)$  that is safe for  $A$ 
4       $A = A \cup \{(u, v)\}$ 
5  return  $A$ 
```

We use the loop invariant as follows:

Initialization: After line 1, the set A trivially satisfies the loop invariant.

Maintenance: The loop in lines 2–4 maintains the invariant by adding only safe edges.

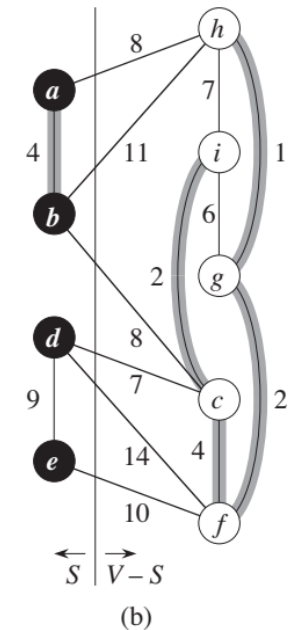
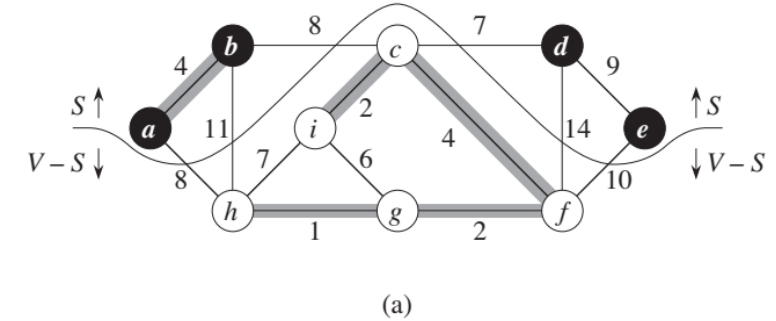
Termination: All edges added to A are in a minimum spanning tree, and so the set A returned in line 5 must be a minimum spanning tree.

Growing a minimum spanning tree

■ Let us consider some definitions:

- A **cut** $(S, V - S)$ of an undirected graph $G = (V, E)$ is a partition of V .
- We say that an edge $(u, v) \in E$ **crosses** the cut if one of its end points is in S and the other is in $V - S$.
- We say that a cut **respects** a set A of edges if no edge in A crosses the cut.
- An edge is a **light edge** crossing a cut if its weight is the minimum of any edge crossing the cut. (able to have more than one light edge)
- The following theorem are the rule for recognizing safe edges.

Let $G = (V, E)$ be a connected, undirected graph with a real-valued weight function w defined on E . Let A be a subset of E that is included in some minimum spanning tree for G , let $(S, V - S)$ be any cut of G that respects A , and let (u, v) be a light edge crossing $(S, V - S)$. Then, edge (u, v) is safe for A .



Complexity Analysis

- The loop in lines 2-4 of GST-MST executes $|V| - 1$ times because it finds one of the $|V| - 1$ edges of a minimum spanning tree in each iteration.
- Initially, when $A = \emptyset$ there are $|V|$ trees in G_A and each iteration reduces that number by 1.

Corollary

Let $G = (V, E)$ be a connected, undirected graph with a real-valued weight function w defined on E . Let A be a subset of E that is included in some minimum spanning tree for G , and let $C = (V_C, E_C)$ be a connected component (tree) in the forest $G_A = (V, A)$. If (u, v) is a light edge connecting C to some other component in G_A , then (u, v) is safe for A .

Kruskal's algorithm

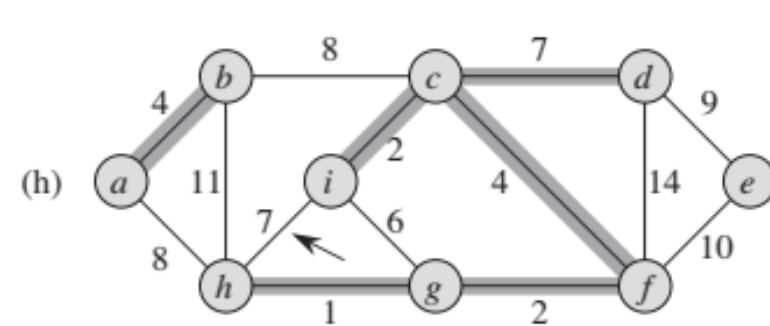
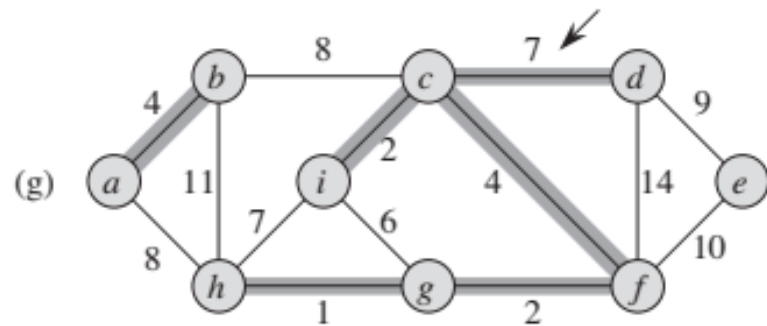
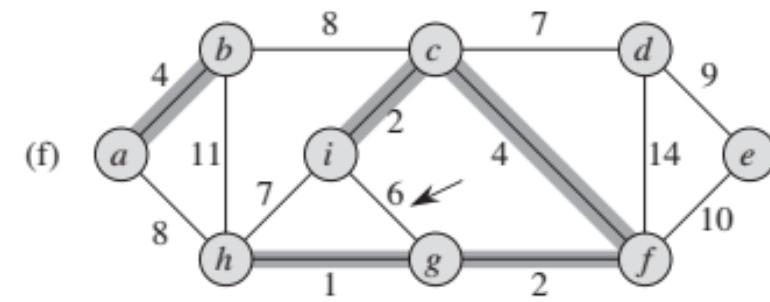
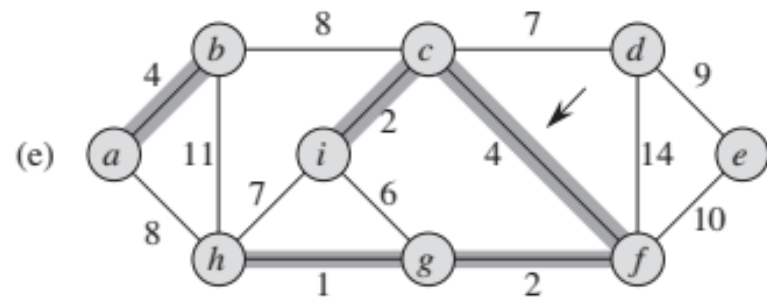
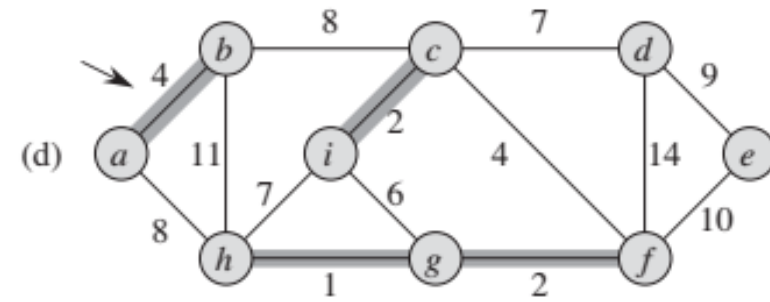
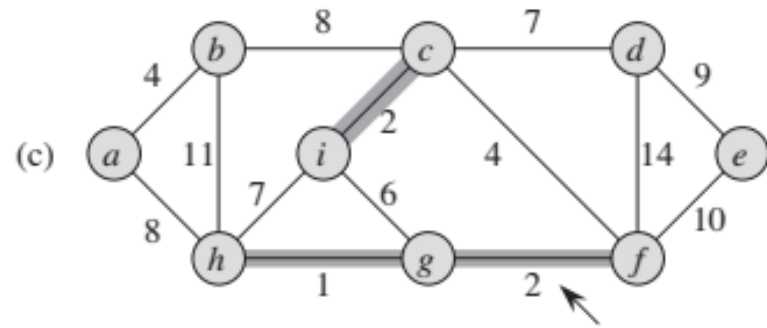
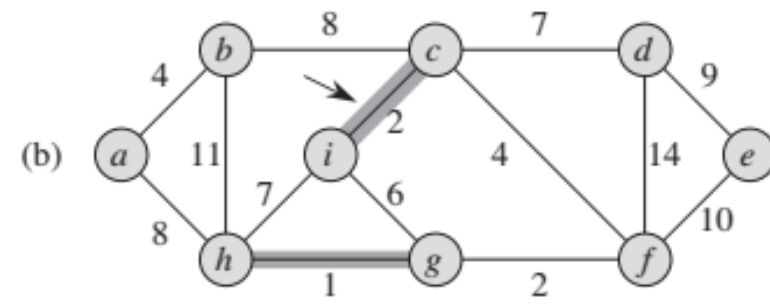
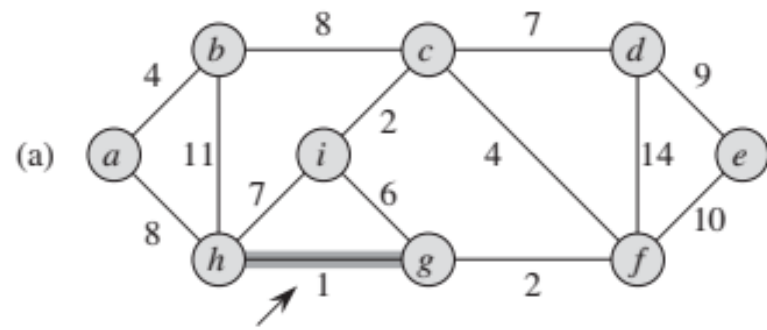
Kruskal's algorithm

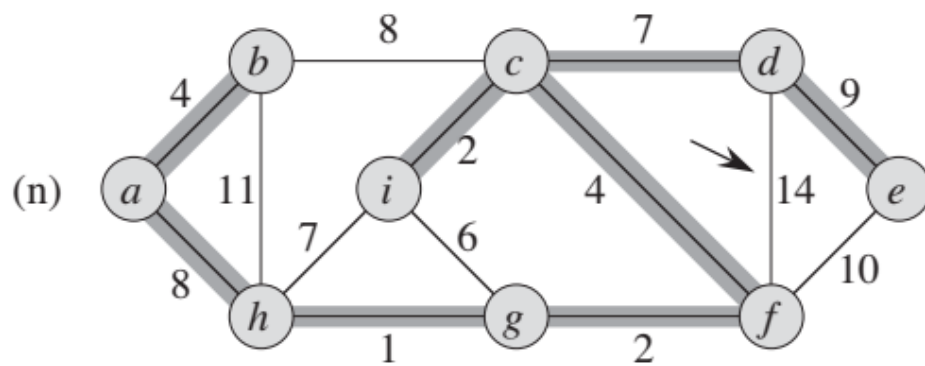
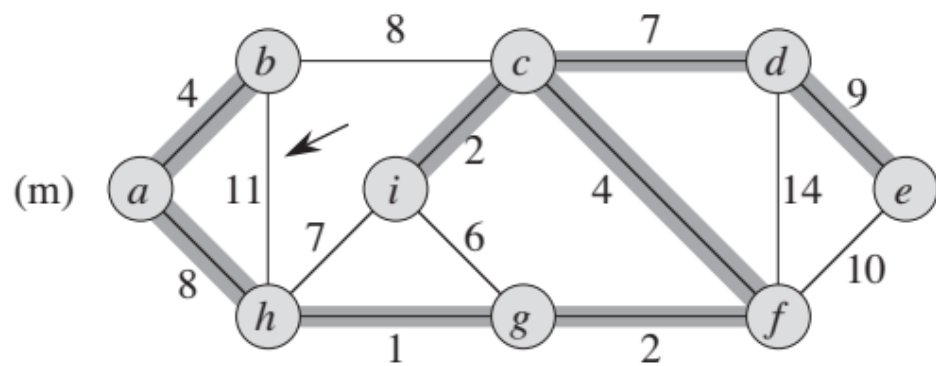
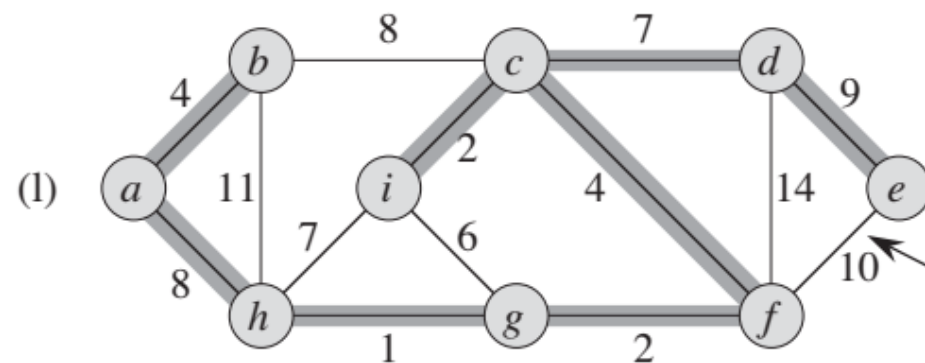
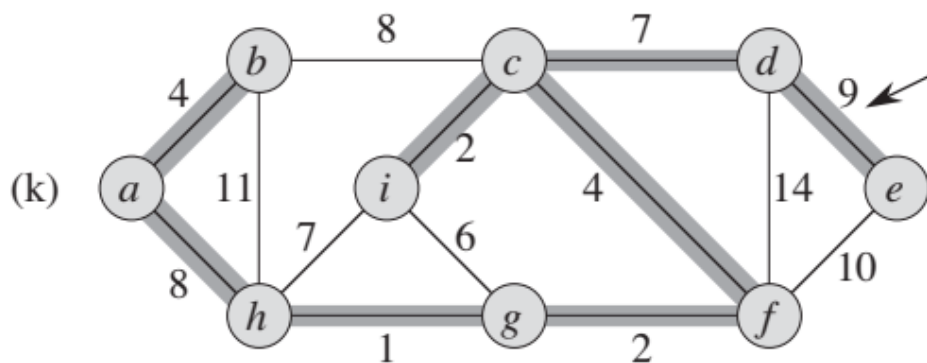
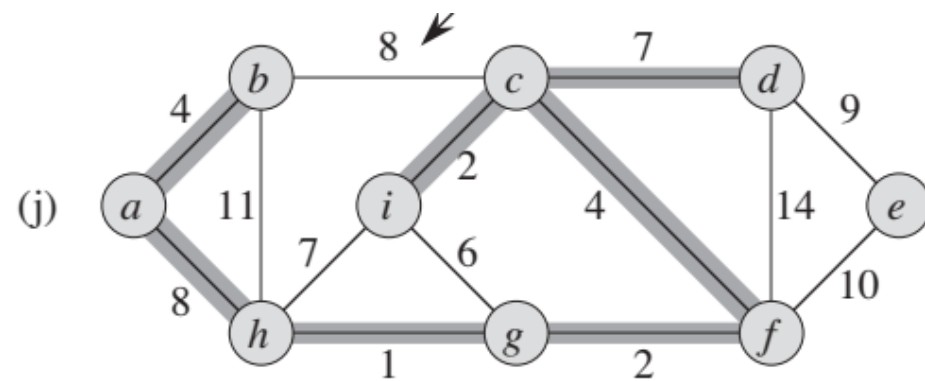
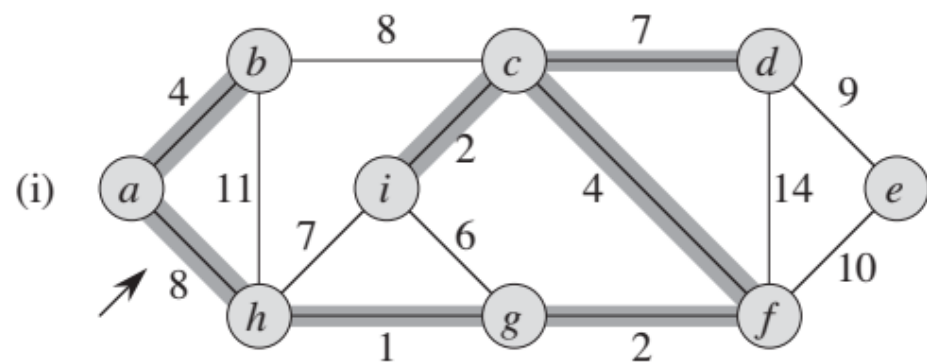
- Kruskal's algorithm finds a safe edge to add to the growing forest by finding, **of all the edges that connect any two trees in the forest**, an edge (u, v) of least weight.
- Let C_1 and C_2 denote the two trees that are connected by (u, v) .
- Since (u, v) must be a **light edge** connecting C_1 to some other tree, the corollary introduced previously (u, v) is a safe edge for C_1
- Kruskal's algorithm qualifies as a **greedy algorithm** because at each step it adds to the forest an edge of least possible weight.

Kruskal's algorithm

MST-KRUSKAL(G, w)

```
1   $A = \emptyset$ 
2  for each vertex  $v \in G.V$ 
3      MAKE-SET( $v$ )
4  sort the edges of  $G.E$  into nondecreasing order by weight  $w$ 
5  for each edge  $(u, v) \in G.E$ , taken in nondecreasing order by weight
6      if FIND-SET( $u$ )  $\neq$  FIND-SET( $v$ )
7           $A = A \cup \{(u, v)\}$ 
8          UNION( $u, v$ )
9  return  $A$ 
```

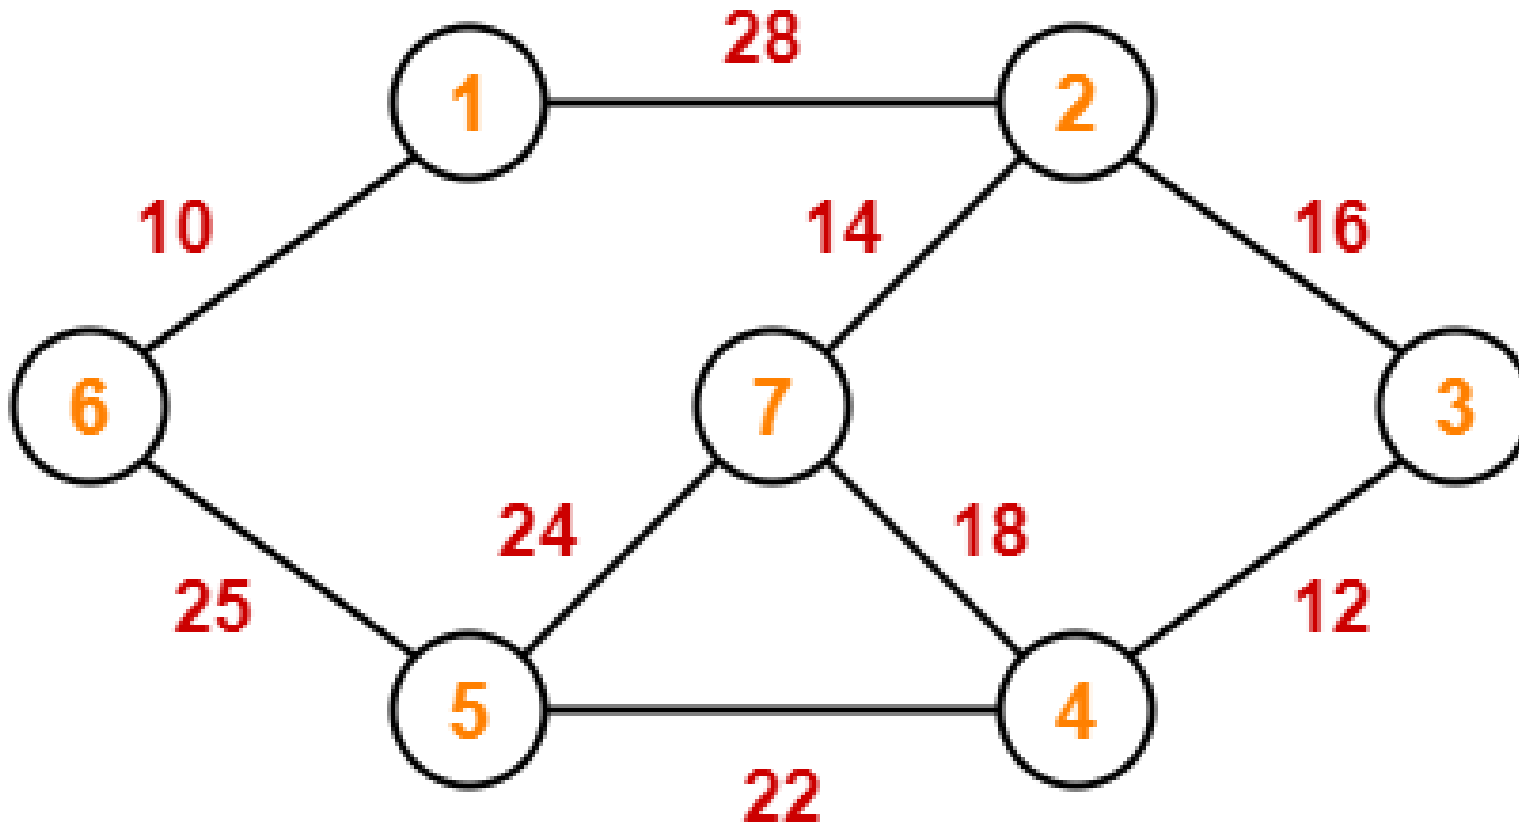


Complexity Analysis

- The running time of Kruskal's algorithm is $O(E \lg V)$.

Quiz#9

Using the Kruskal's algorithm to find
The minimum spanning tree.



Prim's algorithm

Prim's algorithm

- Like Kruskal's algorithm, Prim's algorithm is a special case of the generic minimum-spanning-tree method from Section the previous slide.
- Prim's algorithm operates much like Dijkstra's algorithm for finding shortest paths in a graph.
- Prim's algorithm has the property that the edges in the set A always form a **single tree**.
- The tree starts from an arbitrary root vertex r and grows until the tree spans all the vertices in V .

Prim's algorithm

MST-PRIM(G, w, r)

```
1  for each  $u \in G.V$ 
2       $u.key = \infty$ 
3       $u.\pi = \text{NIL}$ 
4   $r.key = 0$ 
5   $Q = G.V$ 
6  while  $Q \neq \emptyset$ 
7       $u = \text{EXTRACT-MIN}(Q)$ 
8      for each  $v \in G.Adj[u]$ 
9          if  $v \in Q$  and  $w(u, v) < v.key$ 
10              $v.\pi = u$ 
11              $v.key = w(u, v)$ 
```

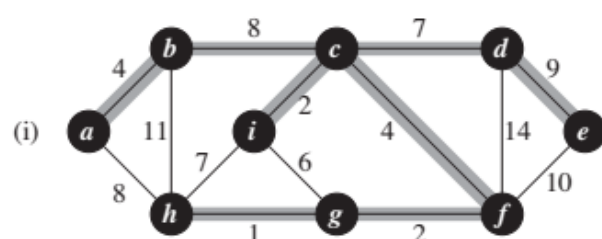
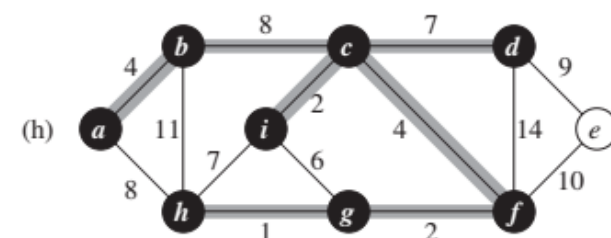
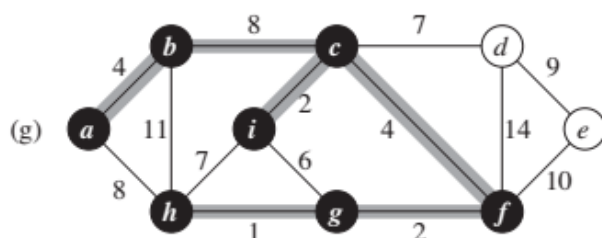
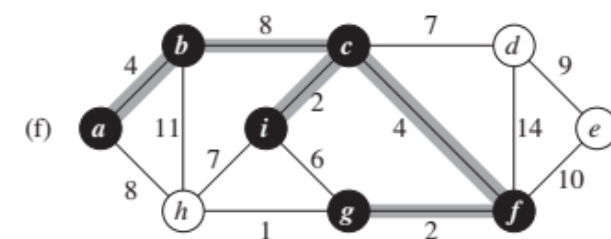
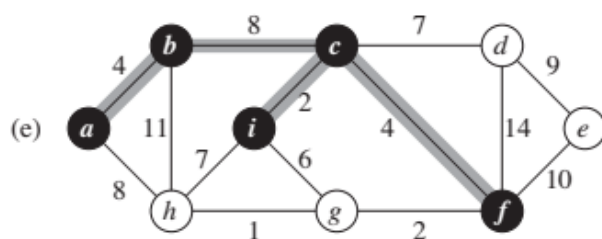
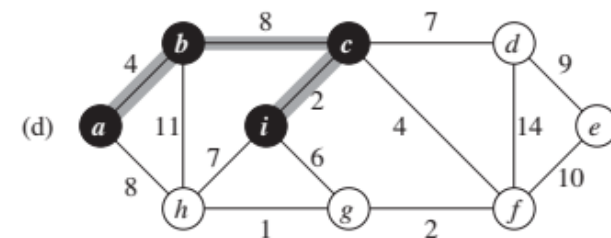
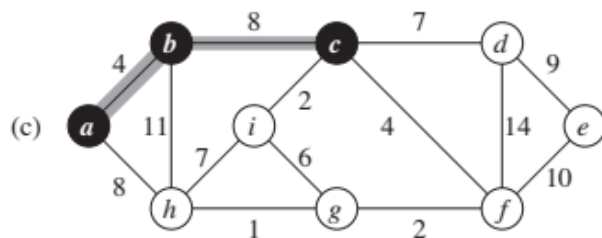
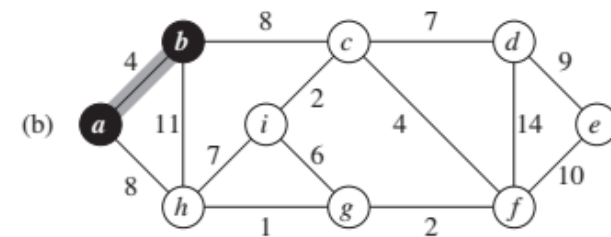
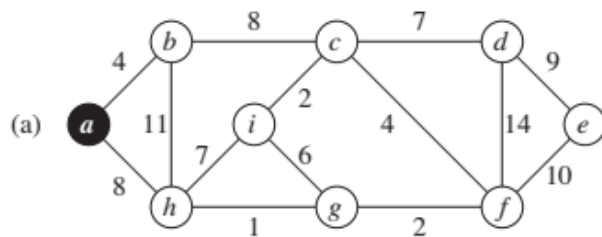
Prior to each iteration of the **while** loop of lines 6–11,

1. $A = \{(v, v.\pi) : v \in V - \{r\} - Q\}$.
2. The vertices already placed into the minimum spanning tree are those in $V - Q$.
3. For all vertices $v \in Q$, if $v.\pi \neq \text{NIL}$, then $v.key < \infty$ and $v.key$ is the weight of a light edge $(v, v.\pi)$ connecting v to some vertex already placed into the minimum spanning tree.

When the algorithm terminates, the min-priority queue Q is empty; the minimum spanning tree A for G is thus

$$A = \{(v, v.\pi) : v \in V - \{r\}\}$$

Example

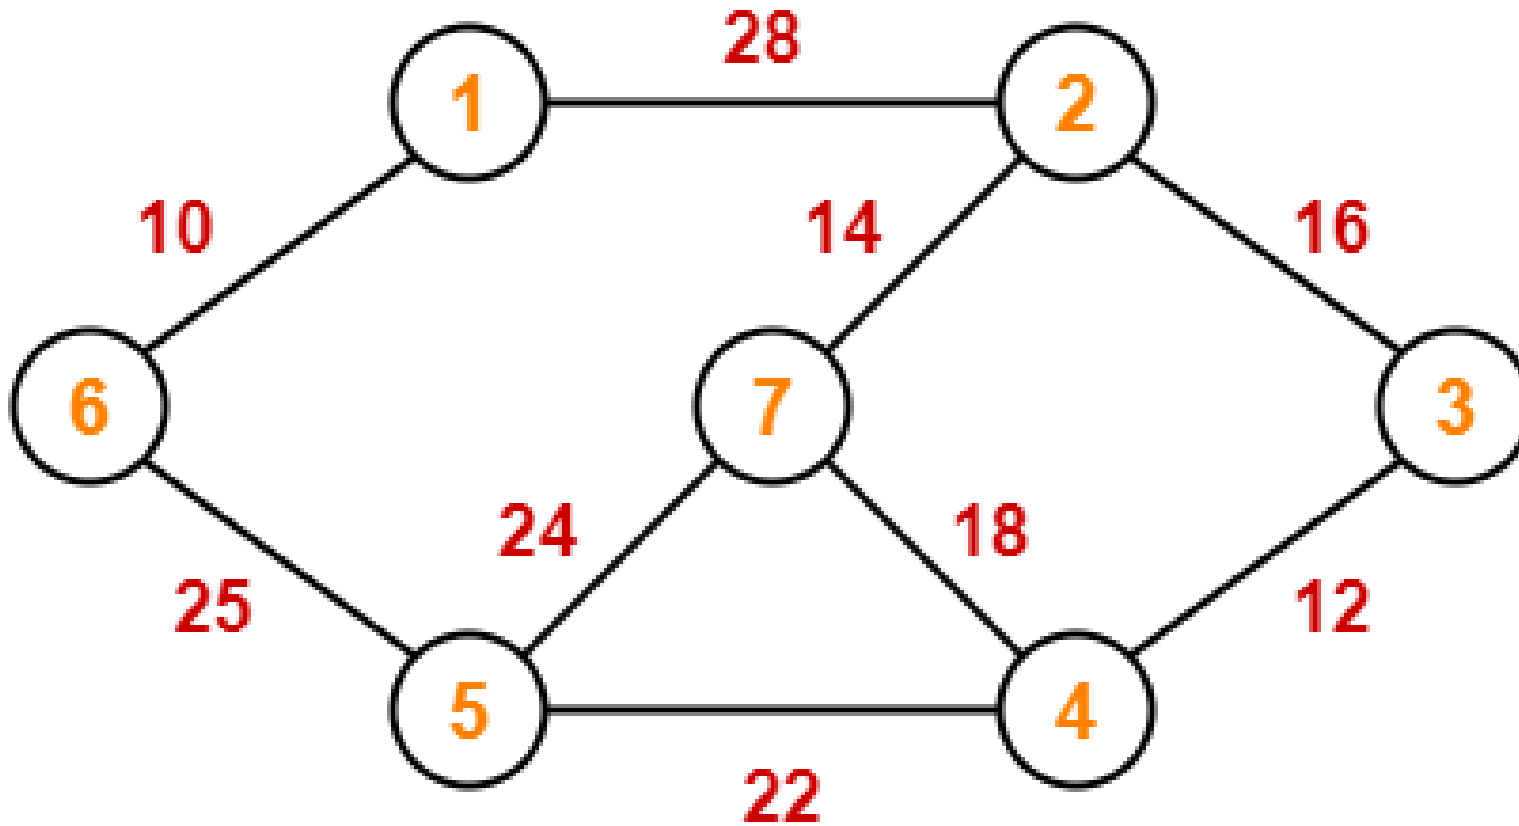


Complexity Analysis

- The running time of Prim's algorithm is $O(E \lg V)$. (Same as Kruskal's algorithm)

Quiz#9

Using the Prim's algorithm to find
The minimum spanning tree.



References